The University of Texas at Austin
Electrical and Computer Engineering
Cockrell School of Engineering

Spring 2022

INTRODUCTION TO COMPUTER VISION

Atlas Wang
Assistant Professor, The University of Texas at Austin

Visual Informatics Group@UT Austin
https://vita-group.github.io/

Many slides here were adapted from **Brown CSCI 1430**

Famous tale in computer vision

- Once, a CMU graduate student asked the famous computer vision scientist **Takeo Kanade**: *"What are the three most important problems in computer vision?"*

- Takeo replied: "**Correspondence, correspondence, correspondence**!"

Visual Correspondence across views

Matching points, patches, edges, or regions across images.

- *Sparse or local correspondence* (picking some “keypoints”)
- *Dense correspondence* (at every pixel)

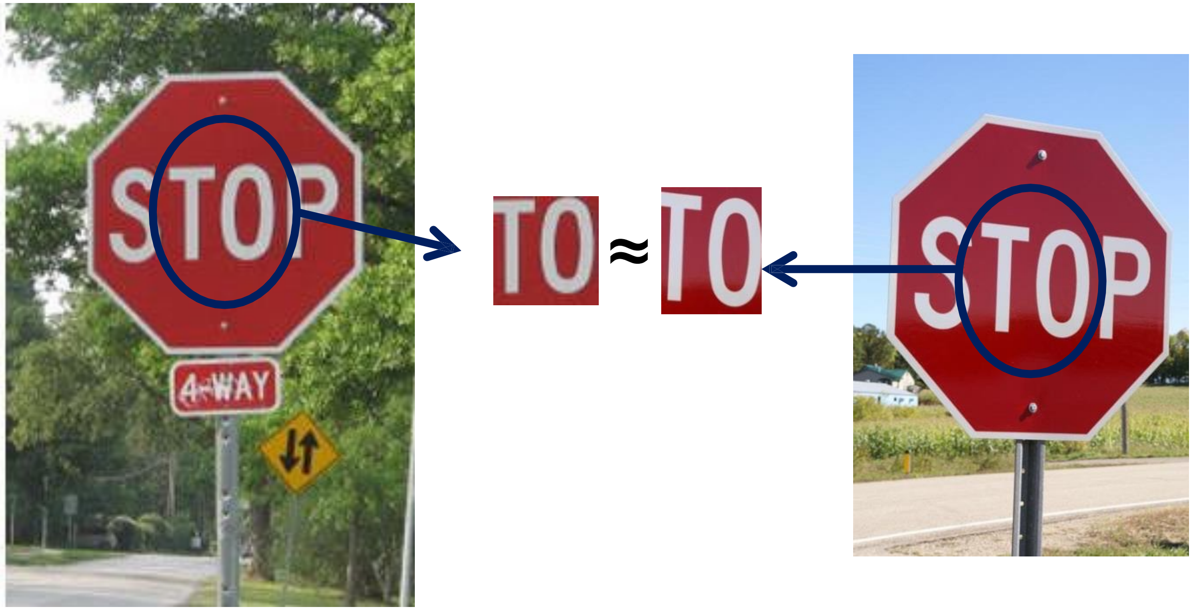

Hays

Fundamental to Applications

- Image alignment
- 3D reconstruction
- Motion tracking (robots, drones, AR)
- Indexing and database retrieval
- Object recognition

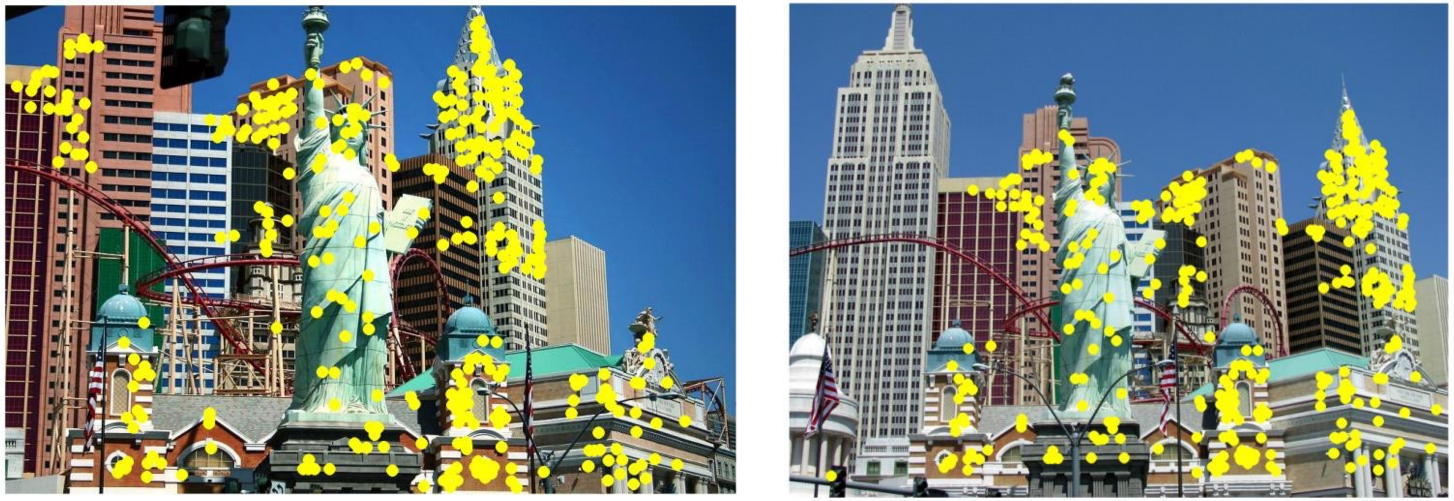

Hays

Example application: Panorama stitching

We have two images –
how do we estimate how to overlay them?

Local features: main components

1) **Detection:**
Find a set of distinctive key points.

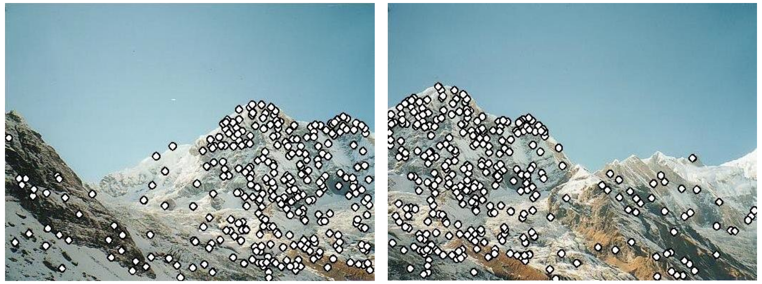

2) **Description:**
Extract feature descriptor around each interest point as vector.

x_1 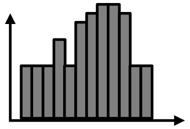

$$\mathbf{x}_1 = [x_1^{(1)}, \ldots, x_d^{(1)}]$$

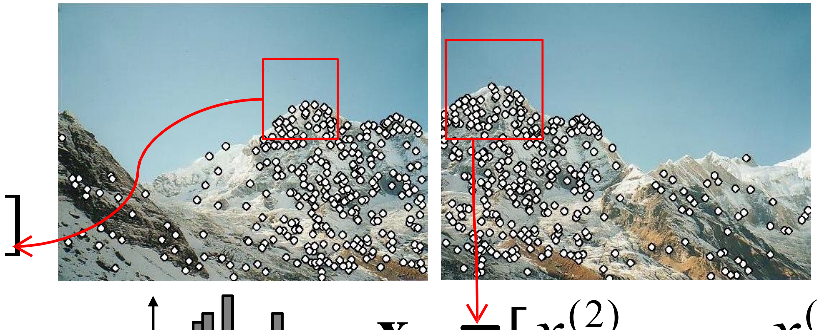

x_2 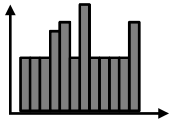

$$\mathbf{x}_2 = [x_1^{(2)}, \ldots, x_d^{(2)}]$$

3) **Matching:**
Compute distance between feature vectors to find correspondence.

$$d(\mathrm{x}_1, \mathrm{x}_2) < T$$

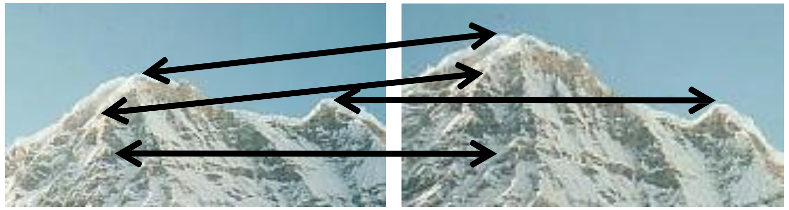

K. Grauman, B. Leibe

Goal: Distinctiveness

We want to be able to reliably determine which point goes with which.

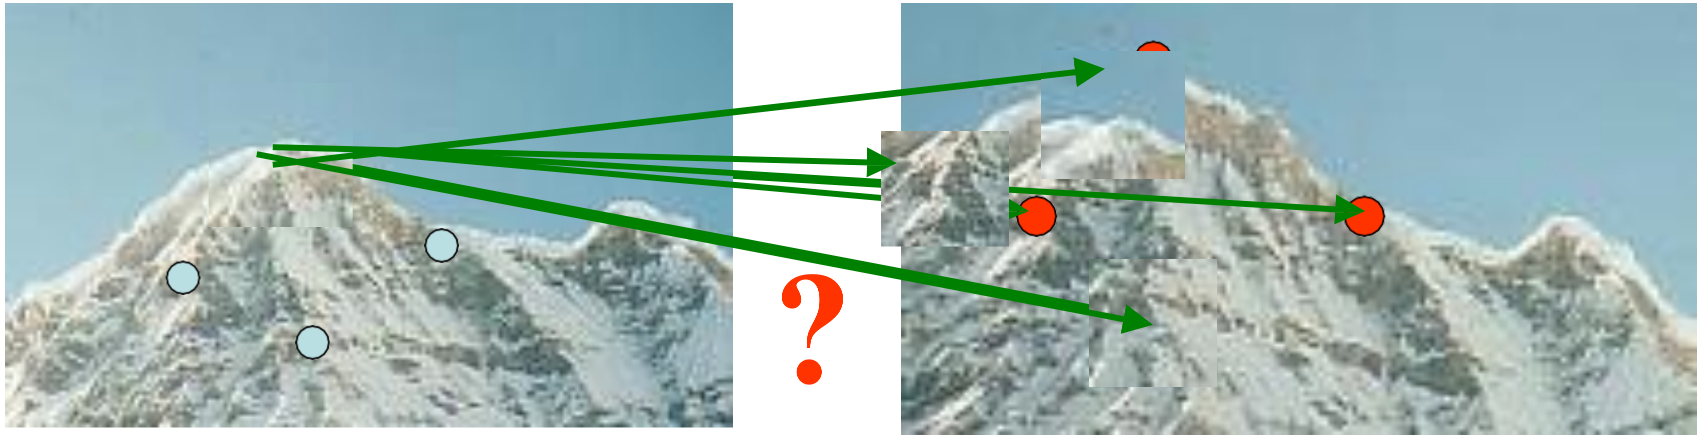

May be difficult in structured environments with repeated elements

Kristen Grauman

Goal: Repeatability

We want to detect (at least some of) the same points in both images.

With these points, there's no chance to find true matches!

Under geometric and photometric variations.

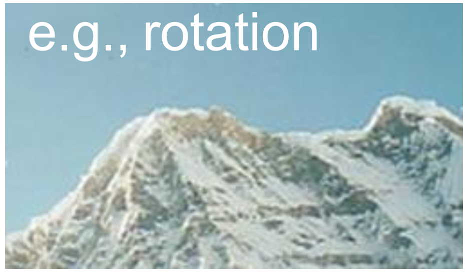

Kristen Grauman

Example: Object Detection

Finding *distinctive* and *repeatable* feature points can be difficult when we want our features to be invariant to large transformations:

- geometric variation (translation, rotation, scale, perspective)
- appearance variation (reflectance, illumination)

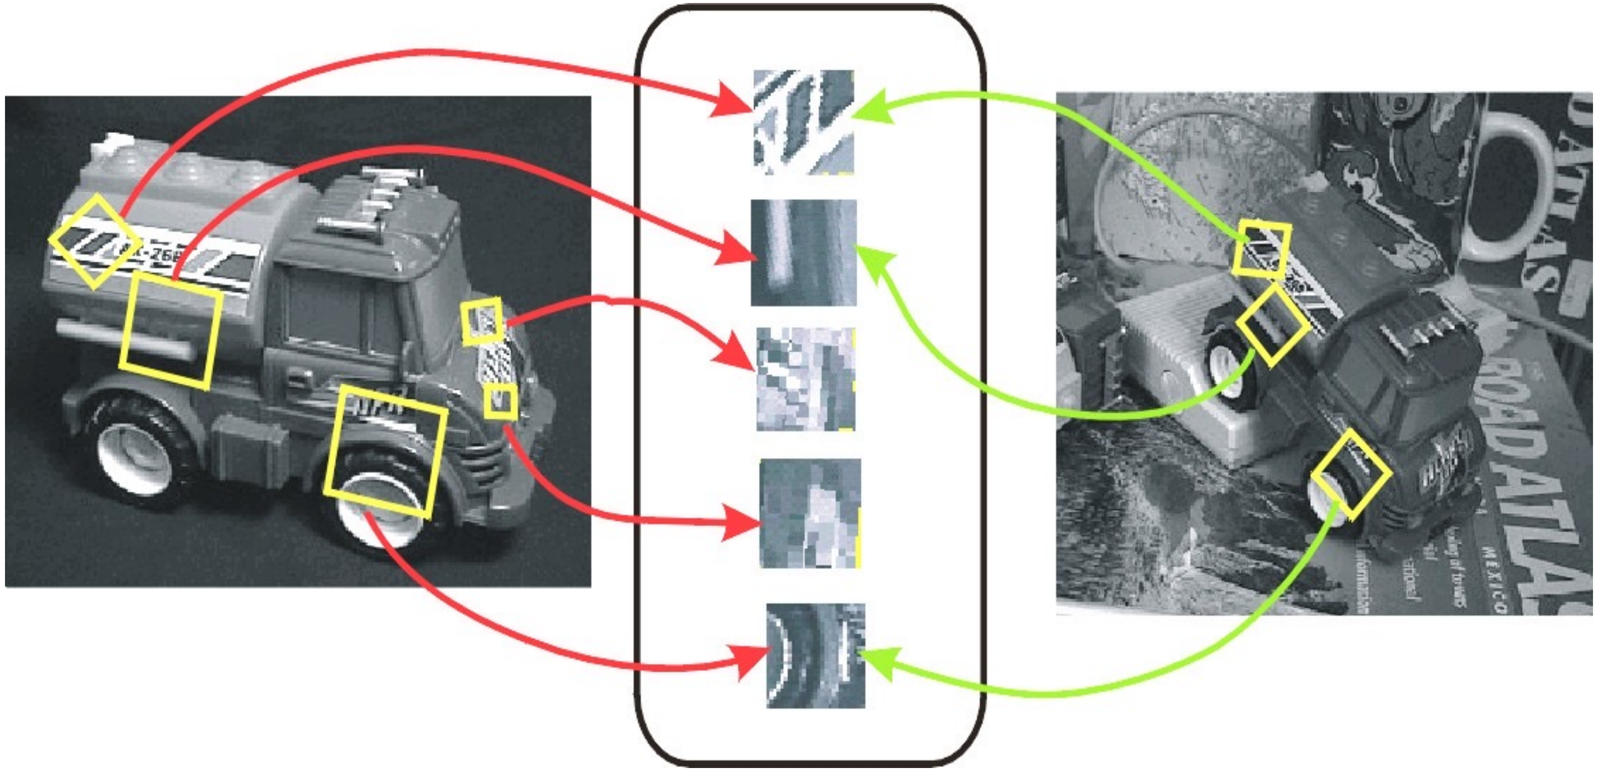

Keypoint Descriptors

James Hays

Goal: Compactness and Efficiency

We want the representation to be as small and as fast as possible

– Much smaller than a whole image

Sometimes, we'd like to run the detection procedure *independently* per image

- Match just the compact descriptors for speed.
- *Difficult!* We don't get to see 'the other image' at match time, e.g., object detection.

Kristen Grauman

Characteristics of good features

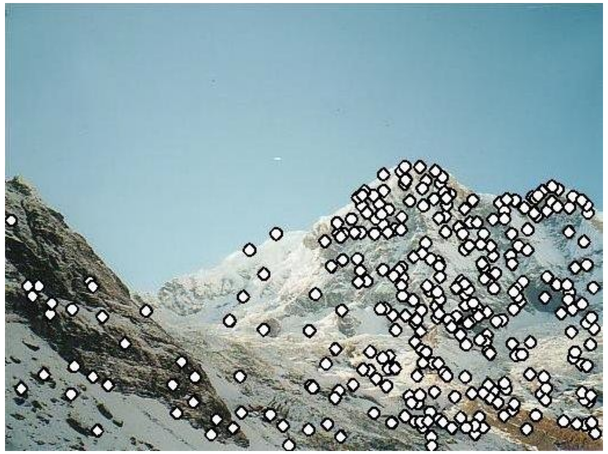

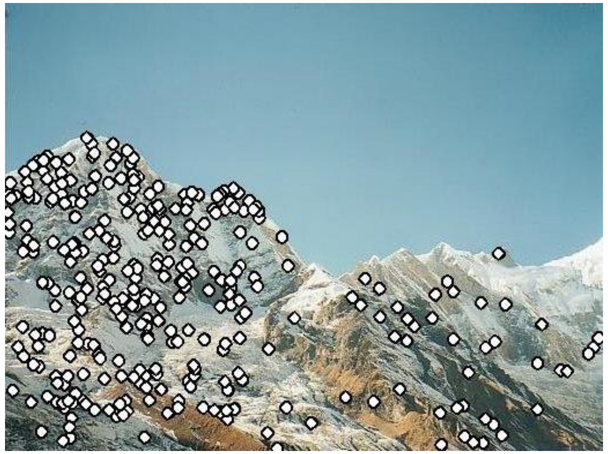

Distinctiveness

Each feature can be uniquely identified

Repeatability

The same feature can be found in several images despite differences:
- geometrically (translation, rotation, scale, perspective)
- photometrically (reflectance, illumination)

Compactness and efficiency

Many fewer features than image pixels; run independently per image

Kristen Grauman

Local features: main components

1) **Detection:**
   Find a set of distinctive key points.

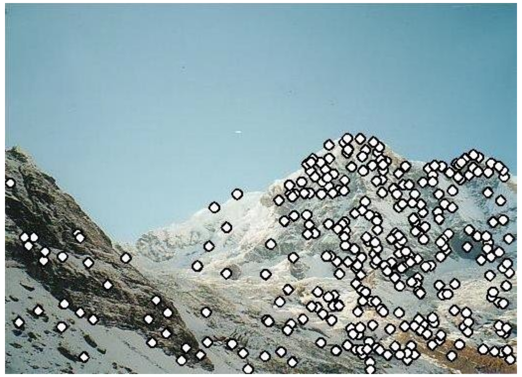

2) **Description:**
   Extract feature descriptor around each interest point as vector.

3) **Matching:**
   Compute distance between feature vectors to find correspondence.

Detection: Basic Idea

We do not know which other image locations the feature will end up being matched against …

But can compute how stable a location is in appearance with respect to small variations in its position

Something that "meaningfully stands out"!

Strategy: Compare image patch against local neighbors

A. Efros

Detection: Basic Idea

Recognize corners by looking at small window.

We want a window shift in *any direction* to give *a large change* in intensity.

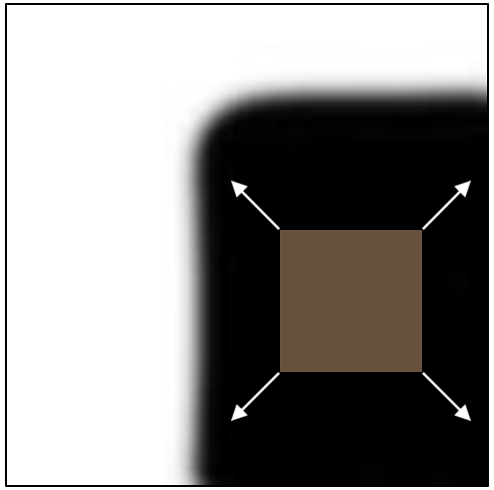

"Flat" region: no change in all directions

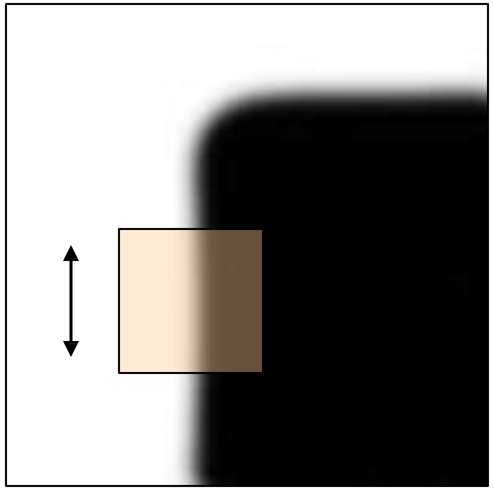

"Edge": no change along the edge direction

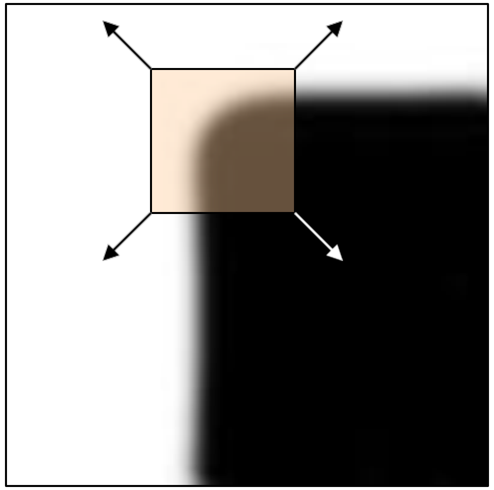

"Corner": significant change in all directions

A. Efros

Finding Corners

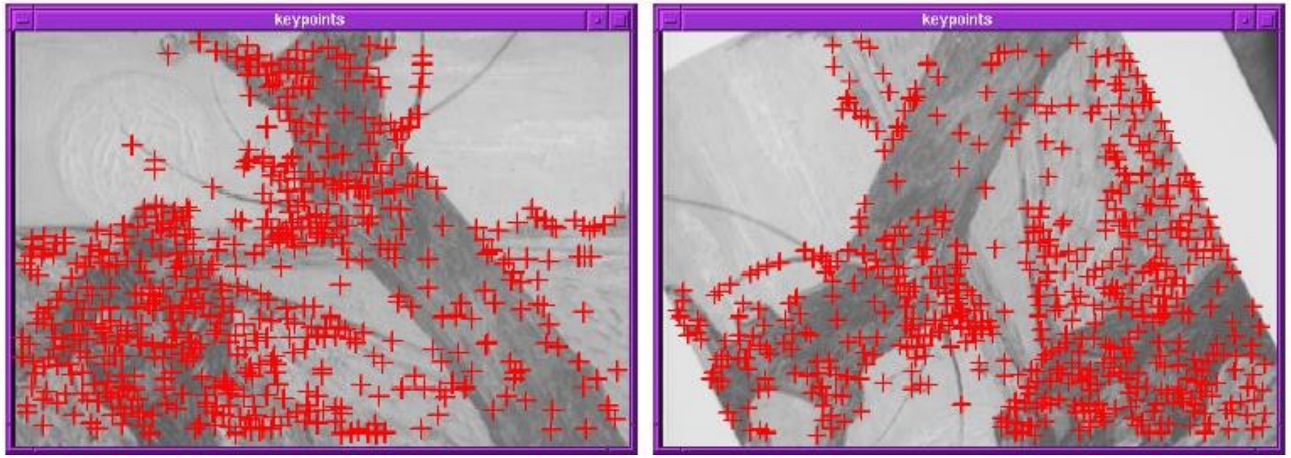

- Key property: in the region around a corner, image gradient has two or more dominant directions
- Corners are repeatable and distinctive

C.Harris and M.Stephens. "A Combined Corner and Edge Detector." *Proceedings of the 4th Alvey Vision Conference*: pages 147--151.

Corner Detection by Auto-correlation

Change in appearance of window $w(x,y)$ for shift $[u,v]$:

$$E(u,v) = \sum_{x,y} w(x,y)\left[I(x+u,y+v) - I(x,y)\right]^2$$

Window function

Shifted intensity

Intensity

Also called 'sum of squared differences'

Window function $w(x,y) =$

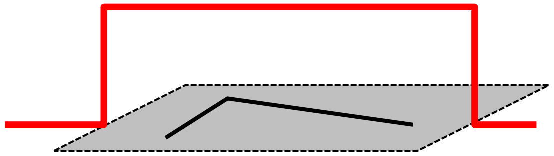

1 in window, 0 outside

or

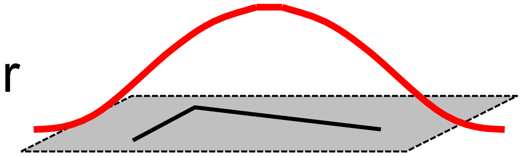

Gaussian

Source: R. Szeliski

Corner Detection by Auto-correlation

Change in appearance of window $w(x,y)$ for shift $[u,v]$:

$$E(u,v) = \sum_{x,y} w(x,y) \left[I(x+u, y+v) - I(x,y) \right]^2$$

$I(x, y)$

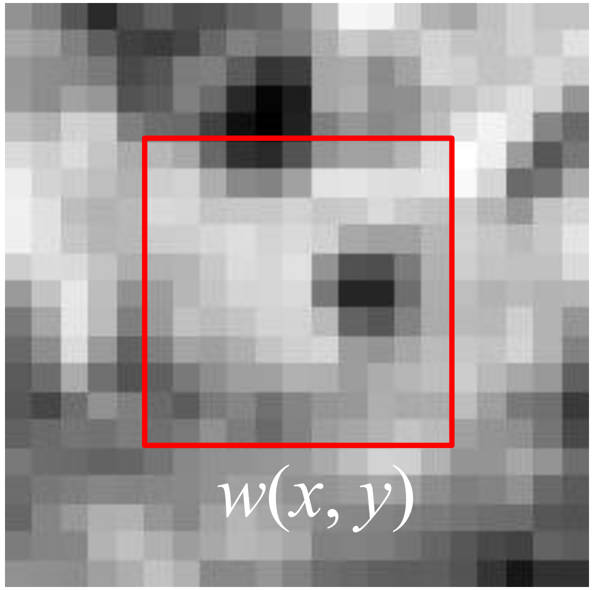

$E(u, v)$

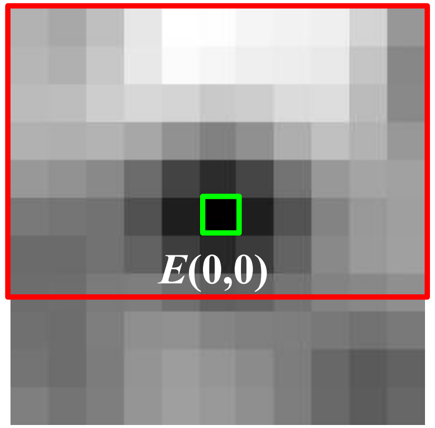

Corner Detection by Auto-correlation

Change in appearance of window $w(x,y)$ for shift $[u,v]$:

$$E(u,v) = \sum_{x,y} w(x,y) \left[I(x+u, y+v) - I(x,y) \right]^2$$

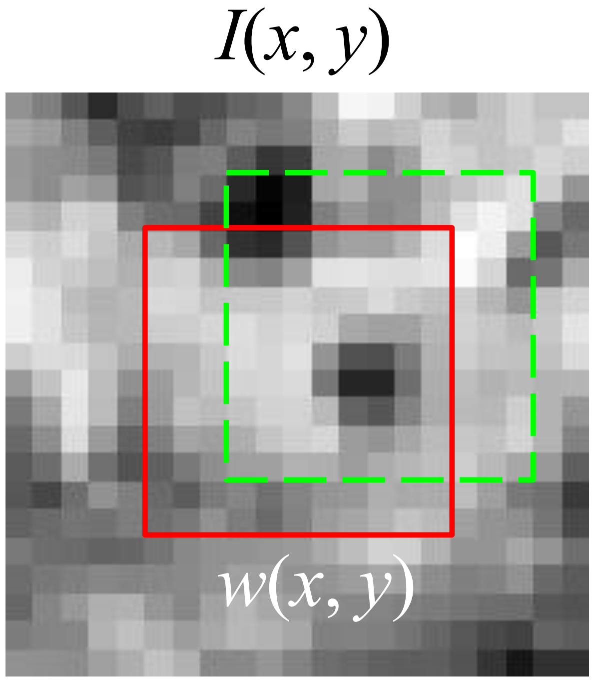

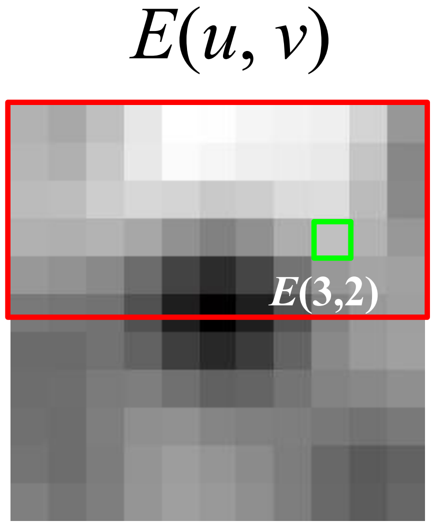

Corner Detection by Auto-correlation

Change in appearance of window $w(x,y)$ for shift $[u,v]$:

$$E(u,v) = \sum_{x,y} w(x,y) \left[I(x+u, y+v) - I(x,y) \right]^2$$

We want to discover how E behaves for small shifts

But this is very slow to compute naively.
O(window_width^2 * shift_range^2 * image_width^2)

O(11^2 * 11^2 * 600^2) = 5.2 billion of these
14.6k ops per image pixel

Corner Detection by Auto-correlation

Change in appearance of window *w*(*x*,*y*) for shift [*u,v*]:

$$E(u,v) = \sum_{x,y} w(x,y)\left[I(x+u,y+v) - I(x,y)\right]^2$$

….But we know the response in *E* that we are looking for – strong peak.

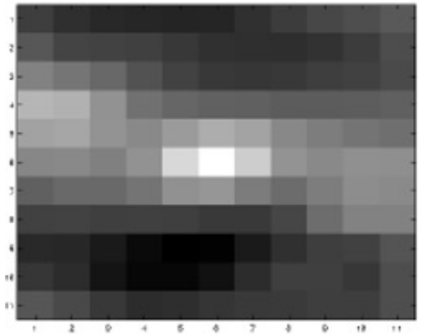

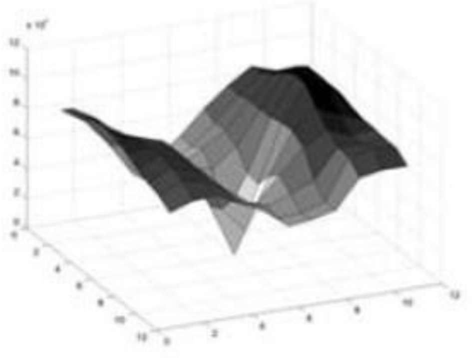

Strategy:

Approximate $E(u,v)$ locally by a quadratic surface, and look for that instead.

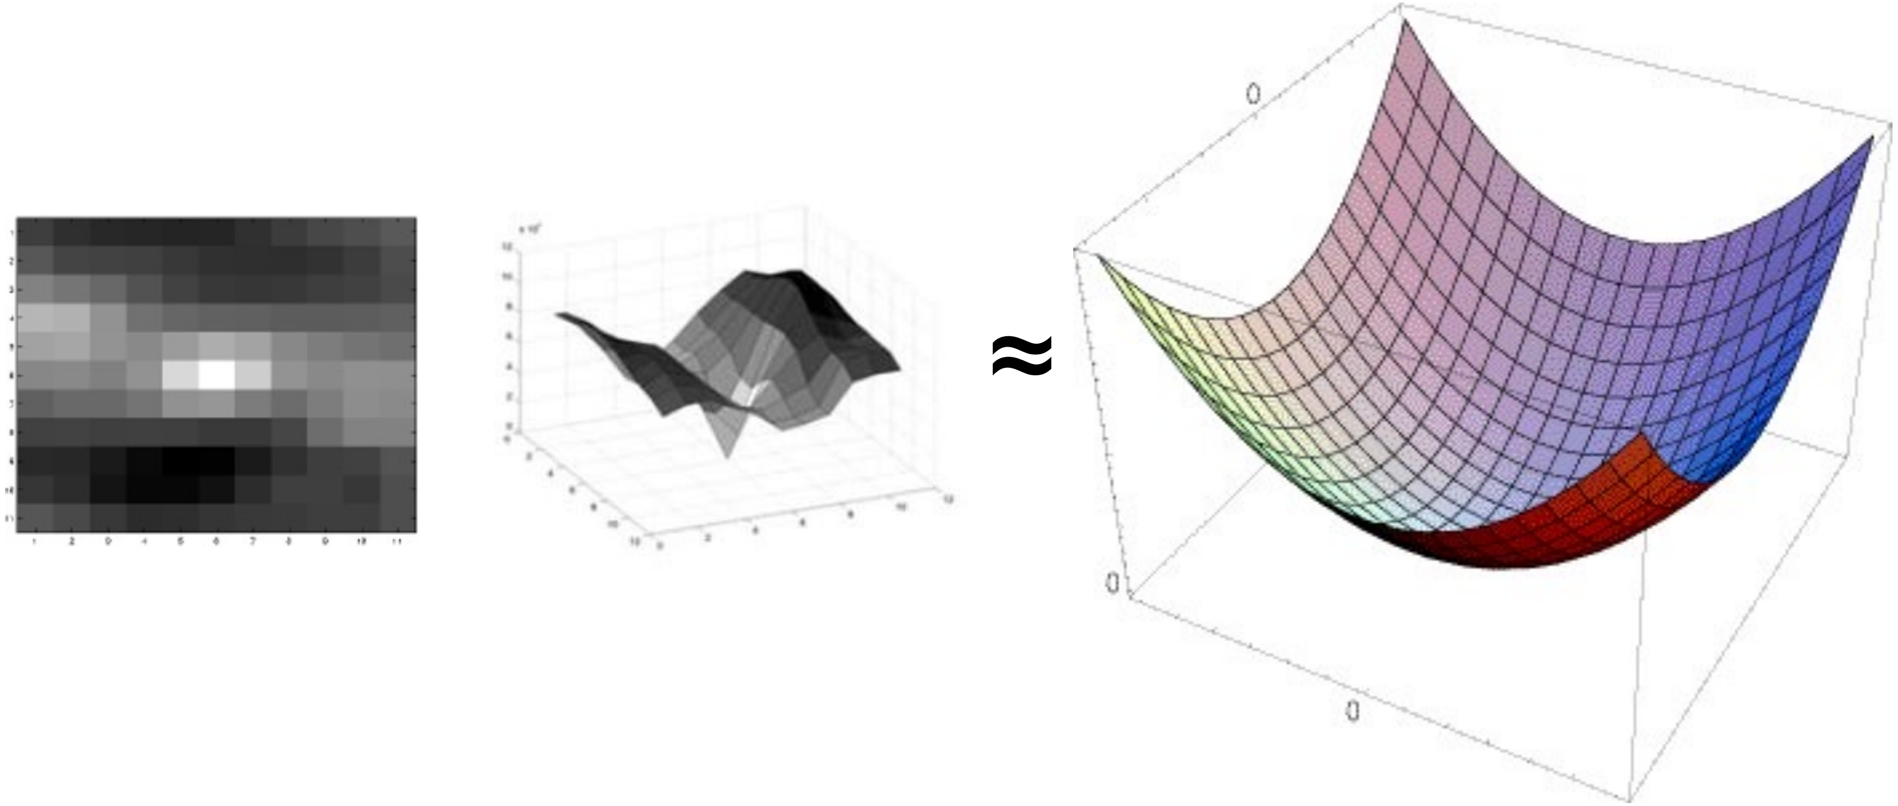

Recall: Taylor series expansion

A function f can be represented by an infinite series of its derivatives at a single point *a*:

$$f(a)+\frac{f'(a)}{1!}(x-a)+\frac{f''(a)}{2!}(x-a)^2+\frac{f'''(a)}{3!}(x-a)^3+\cdots.$$

As we care about window centered, we set *a = 0* (MacLaurin series)

Wikipedia

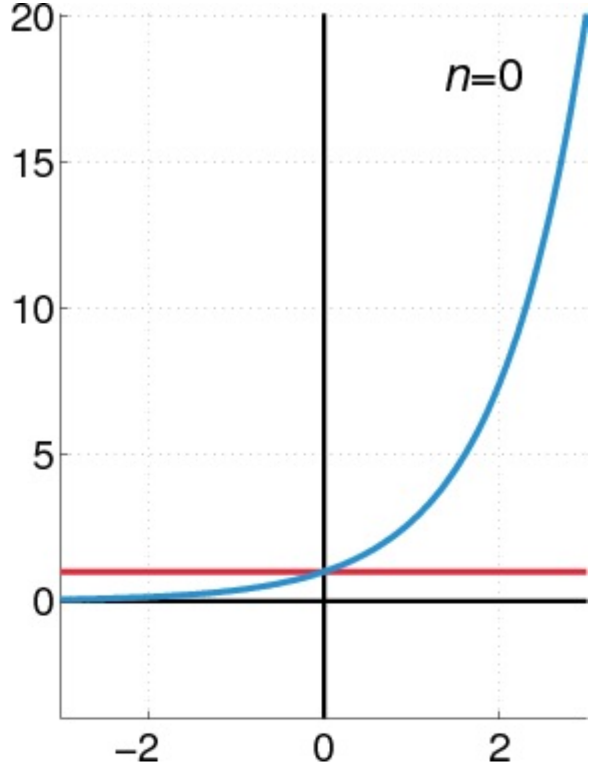

Approximation of $f(x) = e^x$ centered at f(0)

Local quadratic approximation of $E(u,v)$ in the neighborhood of (0,0) is given by the *second-order Taylor expansion*:

$$E(u,v) \approx E(0,0) + [u \quad v]\begin{bmatrix} E_u(0,0) \\ E_v(0,0) \end{bmatrix} + \frac{1}{2}[u \quad v]\begin{bmatrix} E_{uu}(0,0) & E_{uv}(0,0) \\ E_{uv}(0,0) & E_{vv}(0,0) \end{bmatrix}\begin{bmatrix} u \\ v \end{bmatrix}$$

Notation: partial derivative

Local quadratic approximation of *E*(*u*,*v*) in the neighborhood of (0,0) is given by the ***second-order Taylor expansion***:

$$E(u,v) \approx \underbrace{E(0,0)}_{\substack{\text{Ignore function}\\ \text{value; set to 0}}} + [u \;\; v]\underbrace{\begin{bmatrix} E_u(0,0) \\ E_v(0,0) \end{bmatrix}}_{\substack{\text{Ignore first}\\ \text{derivative,}\\ \text{set to 0}}} + \frac{1}{2}[u \;\; v]\underbrace{\begin{bmatrix} E_{uu}(0,0) & E_{uv}(0,0) \\ E_{uv}(0,0) & E_{vv}(0,0) \end{bmatrix}}_{\substack{\text{Just look at shape of}\\ \text{second derivative}\\ \text{(2D quadratic surface)}}}\begin{bmatrix} u \\ v \end{bmatrix}$$

Corner Detection: Mathematics

The quadratic approximation simplifies to

$$E(u,v) \approx [u \;\; v] \begin{bmatrix} E_{uu}(0,0) & E_{uv}(0,0) \\ E_{uv}(0,0) & E_{vv}(0,0) \end{bmatrix} \begin{bmatrix} u \\ v \end{bmatrix}$$

$$\boxed{E(u,v) \approx [u \;\; v] \; M \begin{bmatrix} u \\ v \end{bmatrix}}$$

where *M* is a *second moment matrix / structure tensor* computed from image derivatives:

$$\boxed{M = \sum_{x,y} w(x,y) \begin{bmatrix} I_x^2 & I_x I_y \\ I_x I_y & I_y^2 \end{bmatrix}}$$

$$M = \begin{bmatrix} \sum I_x I_x & \sum I_x I_y \\ \sum I_x I_y & \sum I_y I_y \end{bmatrix} = \sum \begin{bmatrix} I_x \\ I_y \end{bmatrix} [I_x \; I_y] = \sum \nabla I (\nabla I)^T$$

Corners as distinctive interest points

$$M = \sum w(x,y) \begin{bmatrix} I_x I_x & I_x I_y \\ I_x I_y & I_y I_y \end{bmatrix}$$

2 x 2 matrix of image derivatives
(averaged in neighborhood of a point)

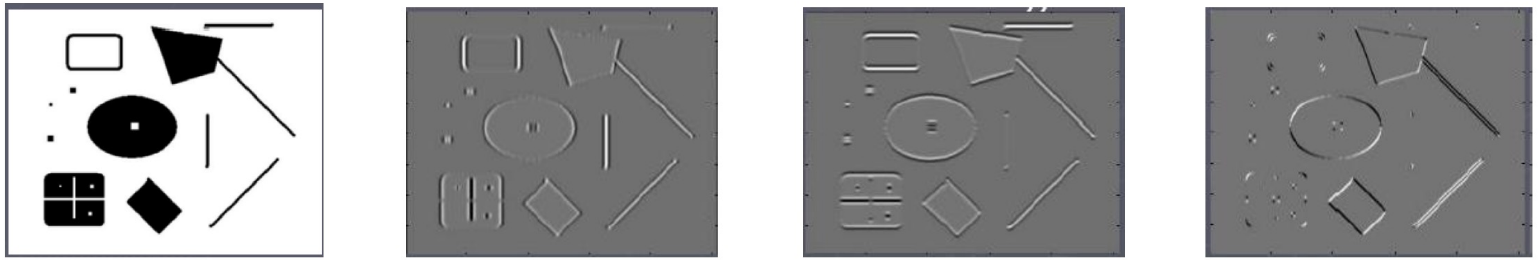

Notation: $I_x \Leftrightarrow \frac{\partial I}{\partial x}$ $I_y \Leftrightarrow \frac{\partial I}{\partial y}$ $I_x I_y \Leftrightarrow \frac{\partial I}{\partial x}\frac{\partial I}{\partial y}$

James Hays

Let’s go back to our goal: corner detection

$$M = \sum_{x,y} w(x,y) \begin{bmatrix} I_x^2 & I_x I_y \\ I_x I_y & I_y^2 \end{bmatrix}$$

- For detecting “cornerness”:
  - Do we care about the change orientation? No
  - Do we care about the change “steepness”? Yes, that is “all we need”
- So, looking at the M approximation now, what we really want?
  - **What if I^2_x, I^2_y, $I_x I_y$ are all small?** No variations -> flat area
  - **What if only I^2_x is large?** Only x-direction has large variations -> Edge
  - **How about only large I^2_y, or $I_x I_y$?** Same thing (edge)
  - **Then, how about letting I^2_x, I^2_y, $I_x I_y$ all be large?**
    - Sufficient, but not necessary…
    - **The missing key: Rotation Invariance**

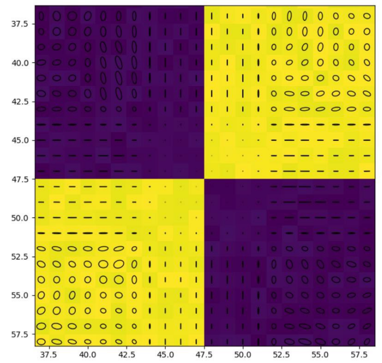

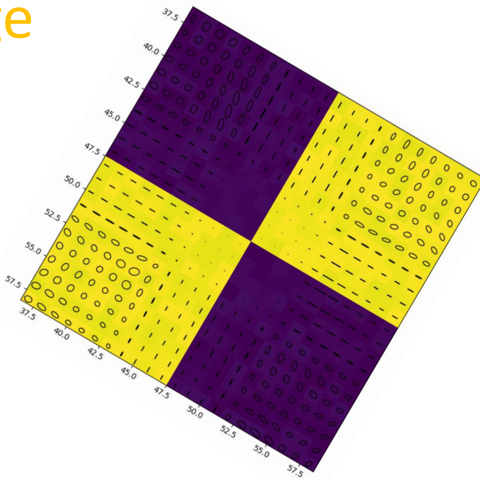

Eigenvalue Analysis (your old friend: PCA)

- **Goal:** Describe the "overall intensity variations" in the window, *regardless of rotation*!

Intensity change in shifting window: eigenvalue analysis

$$E(u,v) \cong [u,v] \quad M \begin{bmatrix} u \\ v \end{bmatrix}$$

λ_1, λ_2 – eigenvalues of M

direction of the fastest change

direction of the slowest change

Ellipse $E(u,v)$ = const

$(\lambda_{max})^{-1/2}$

$(\lambda_{min})^{-1/2}$

Categorizing image points using *M* eigenvalues

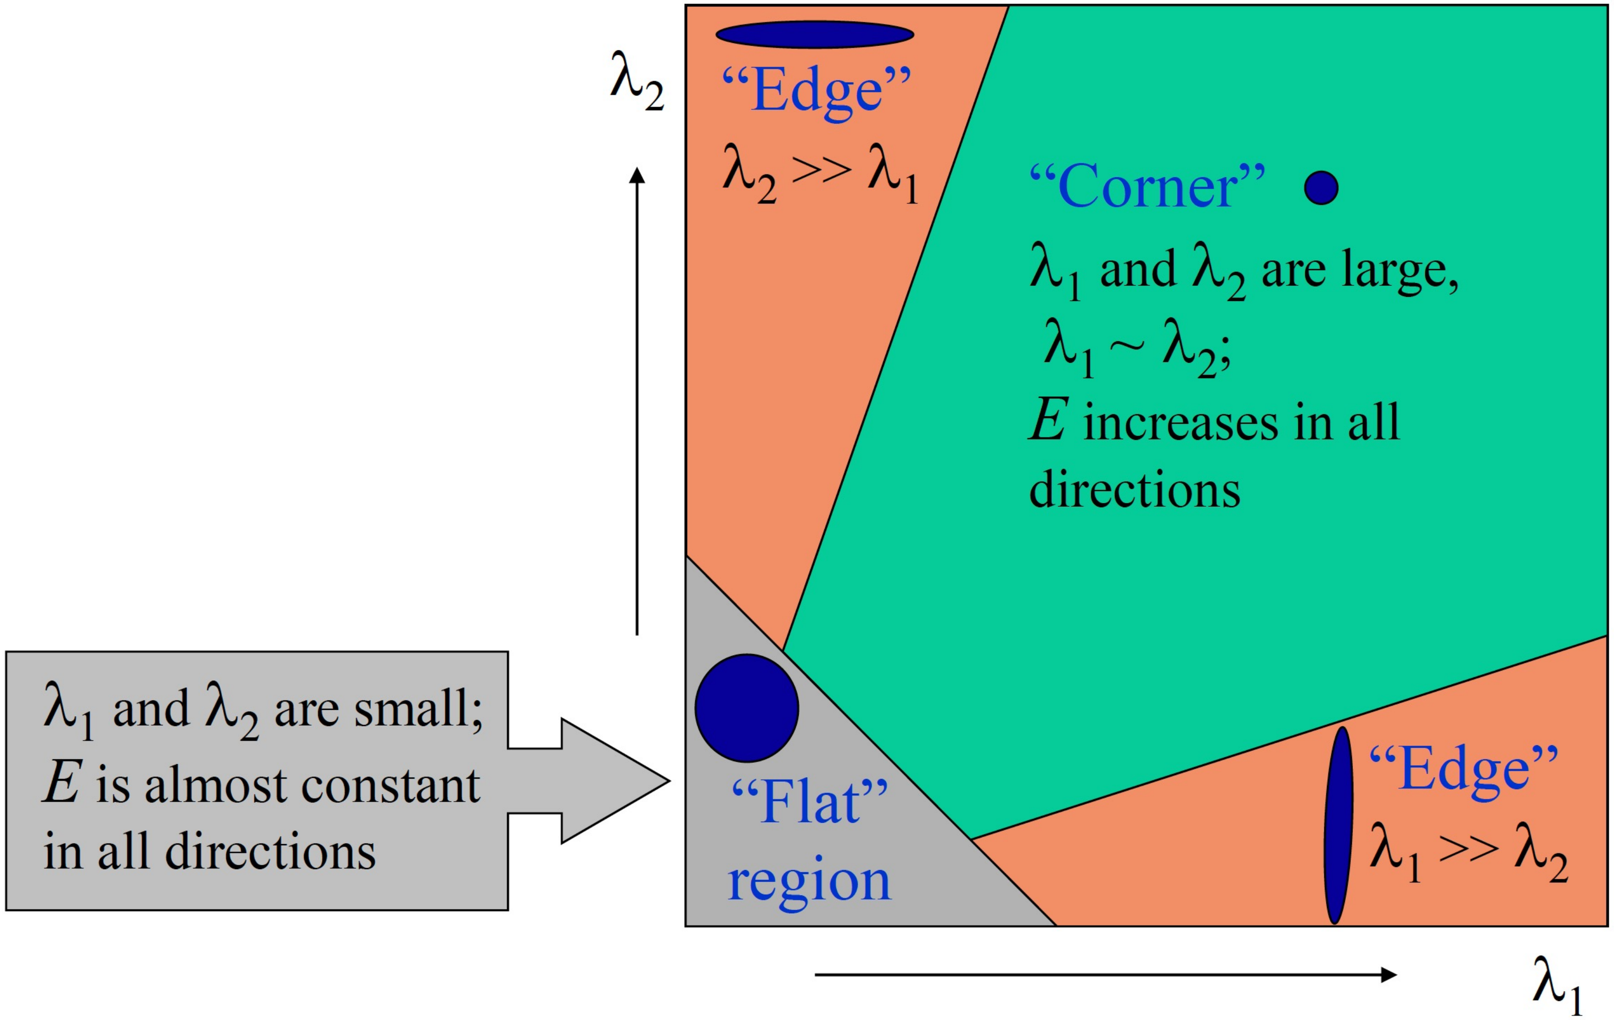

Categorizing image points using *M* eigenvalues

Cornerness score:

$$C = \lambda_1\lambda_2 - \alpha(\lambda_1 + \lambda_2)^2$$

α: some constant (~0.04 to 0.06)

Could You Sense Why?

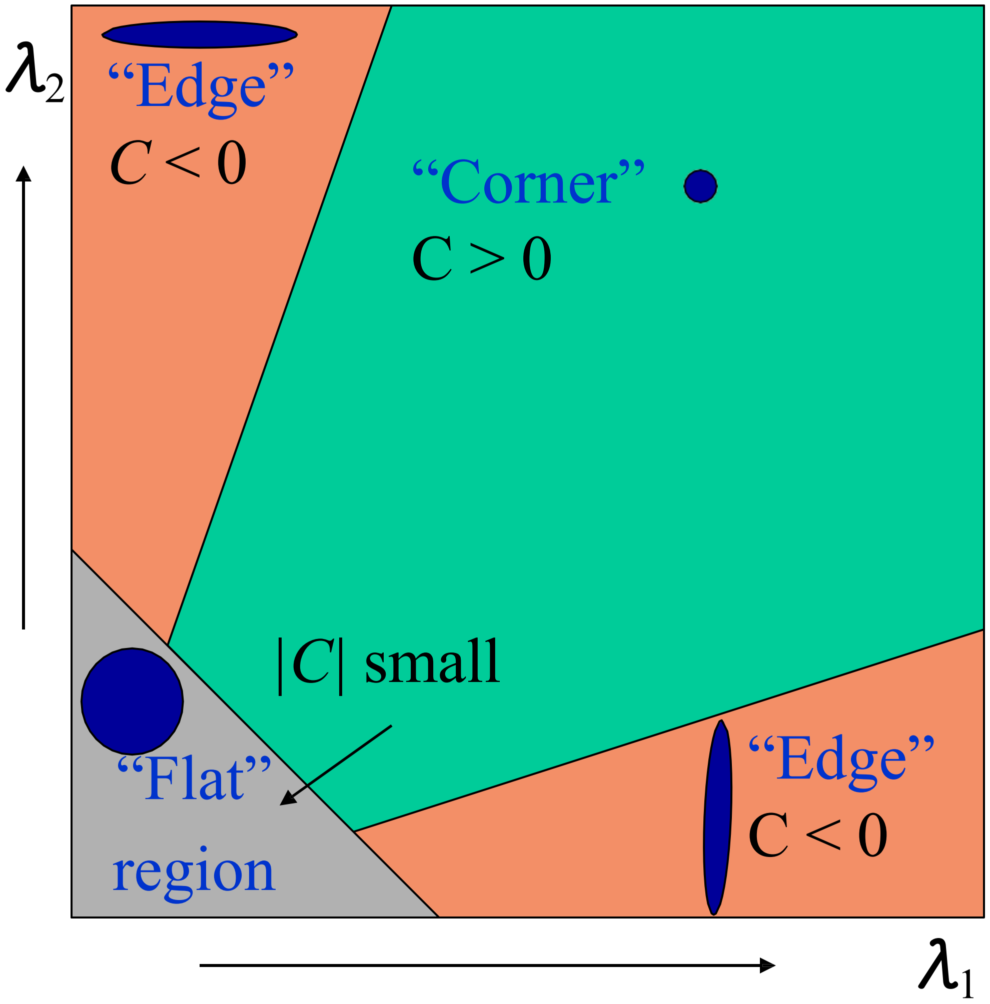

Categorizing image points using *M* eigenvalues

Cornerness score:

$$C = \lambda_1 \lambda_2 - \alpha(\lambda_1 + \lambda_2)^2$$

α: some constant (~0.04 to 0.06)

Remember your linear algebra:

Determinant: $\det(A) = \prod_{i=1}^{n} \lambda_i = \lambda_1 \lambda_2 \cdots \lambda_n$.
(diagonal matrices)

Trace: $\operatorname{tr}(A) = \sum_i \lambda_i$.

$$C = \det(M) - \alpha \operatorname{Tr}(M)$$

Avoids explicit eigenvalue computation!

(many fast algorithms to directly estimate det/Tr)

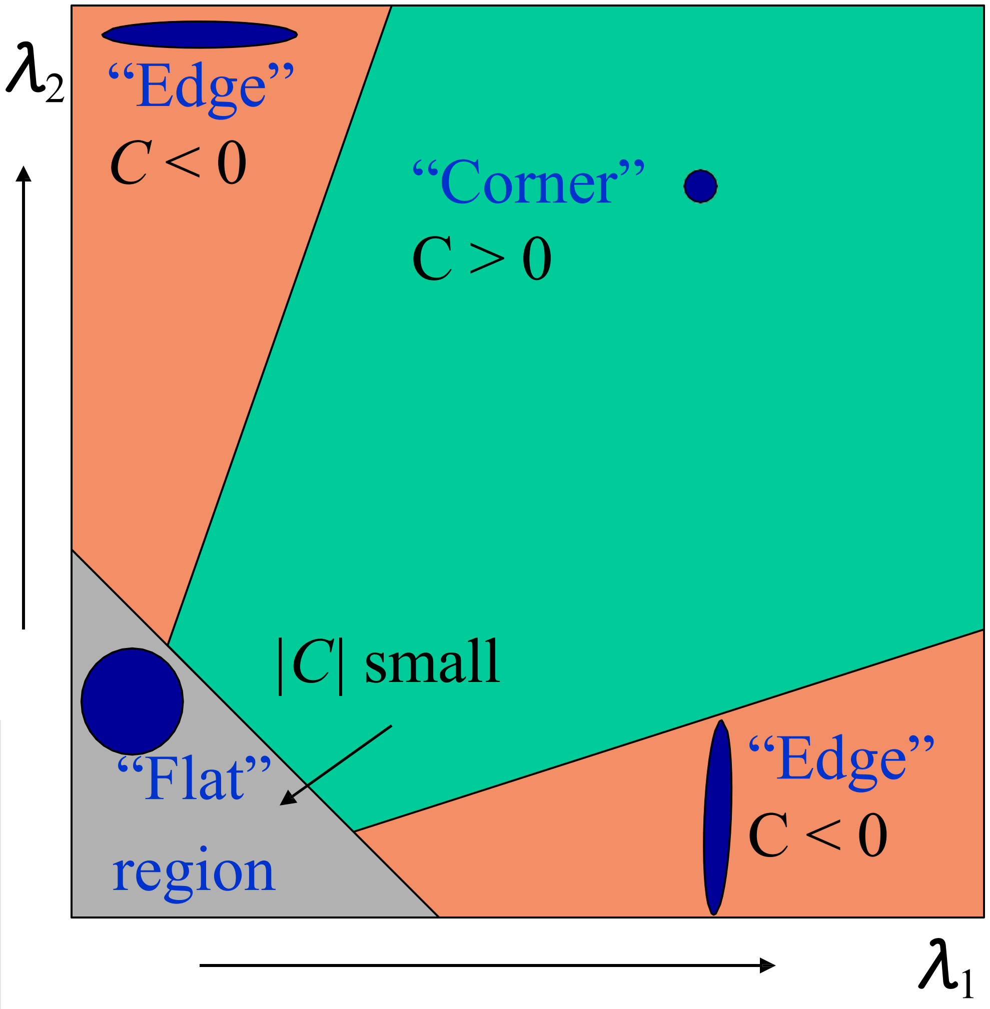

This is the famous Harris corner detector!

1) Compute *M* matrix for each window to recover a *cornerness* score C.

   Note: We can find *M* purely from the per-pixel image derivatives!

2) Threshold to find pixels which give large corner response (C > threshold).

3) Find the local maxima pixels, i.e., non-maximal suppression.

C.Harris and M.Stephens. "A Combined Corner and Edge Detector." *Proceedings of the 4th Alvey Vision Conference*: pages 147—151, 1988.

Harris Corner Detector [Harris88]

James Hays

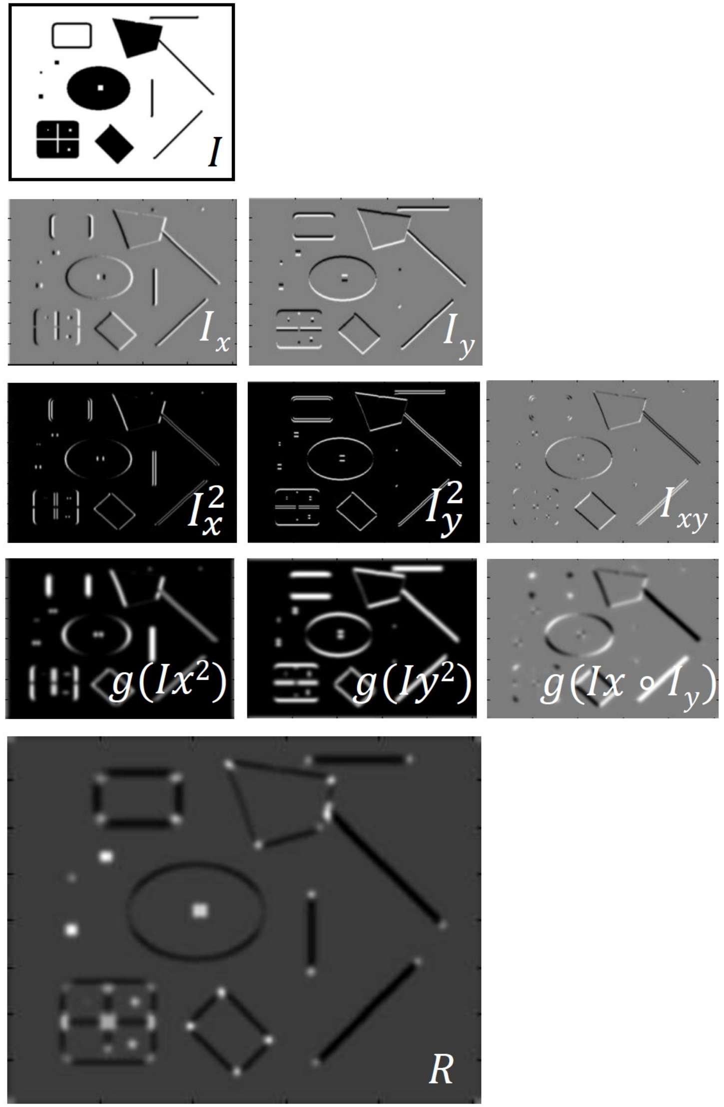

0. Input image
   We want to compute M at each pixel.

1. Compute image derivatives (optionally, blur first).

2. Compute M components
   as squares of derivatives.

3. Gaussian filter *g()* with width σ

   $$= g(I_x^2),\ g(I_y^2),\ g(I_x \circ I_y)$$

Reminder: $a \circ b$ is Hadamard product (element-wise multiplication)

4. Compute cornerness

   $$C = \det(M) - \alpha\ \mathrm{trace}(M)^2$$
   $$= g(I_x^2) \circ g(I_y^2) - g(I_x \circ I_y)^2$$
   $$-\alpha\left[g(I_x^2) + g(I_y^2)\right]^2$$

5. Threshold on C to pick high cornerness

6. Non-maximal suppression to pick peaks.

Harris Detector: Steps

Harris Detector: Steps

Compute corner response C

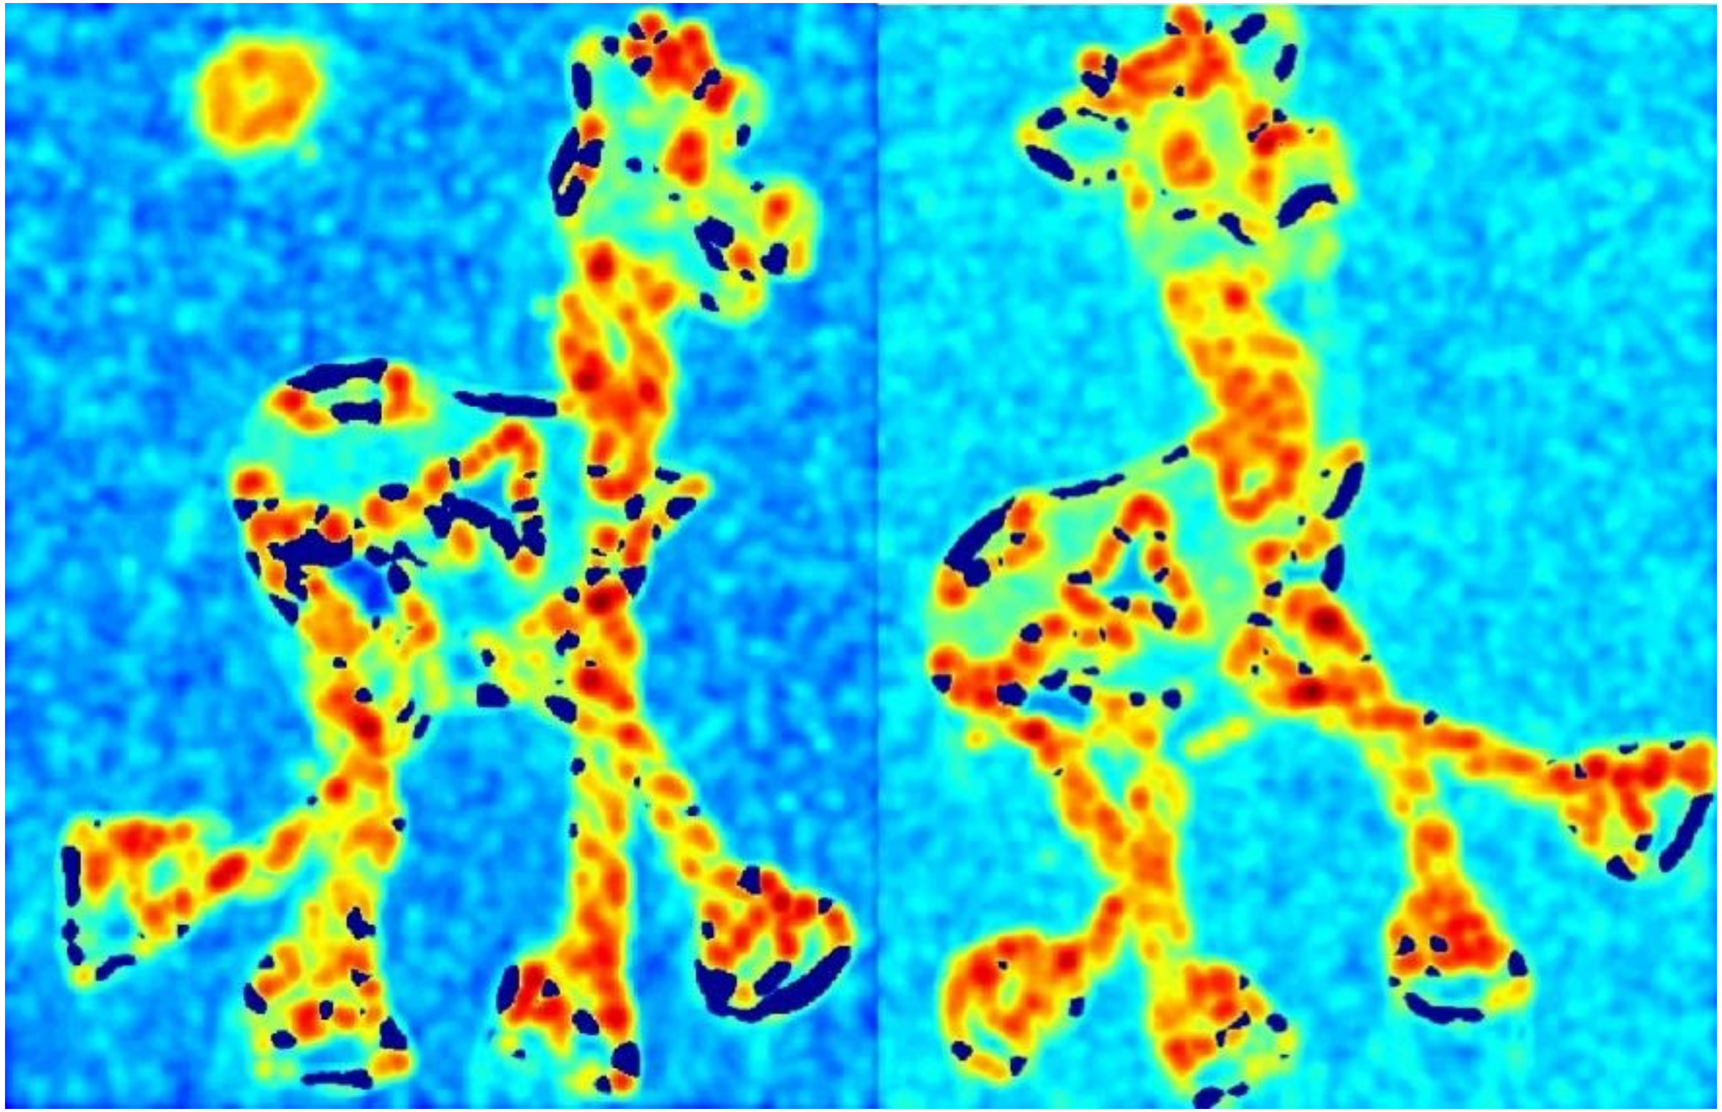

Harris Detector: Steps

Find points with large corner response: C > threshold

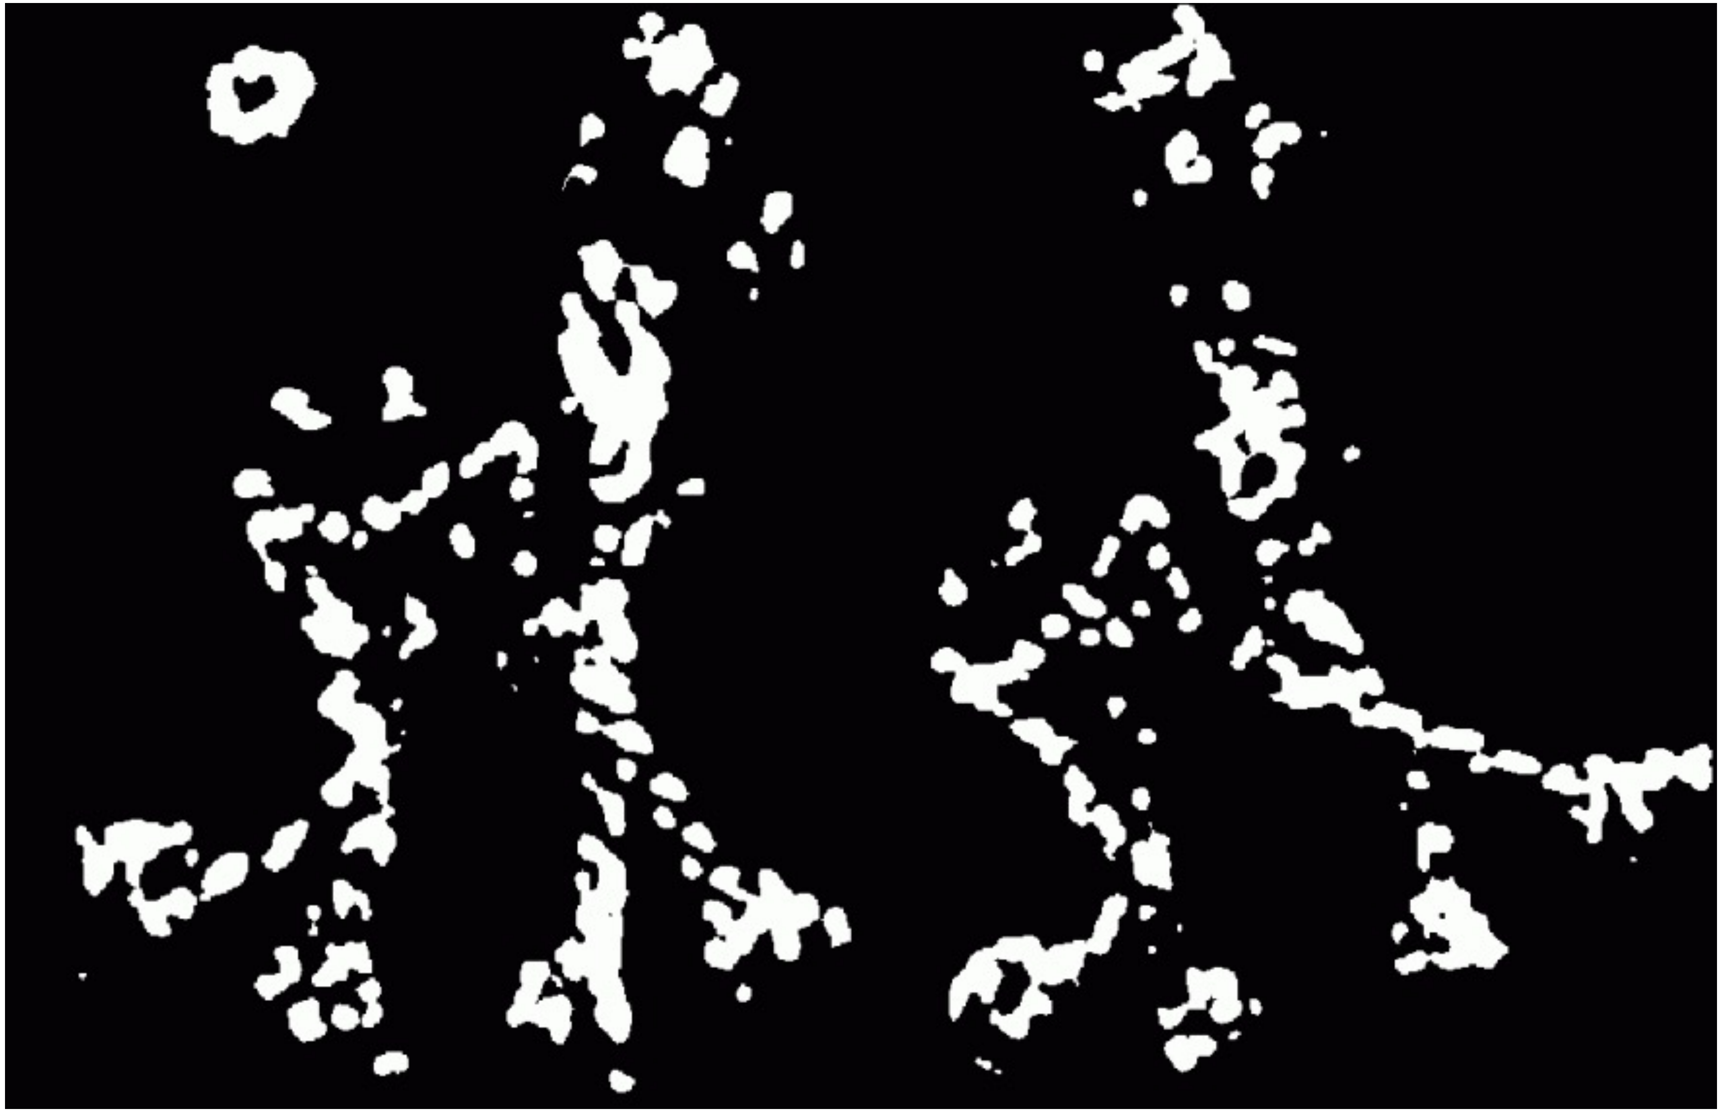

Harris Detector: Steps

Take only the points of local maxima of C

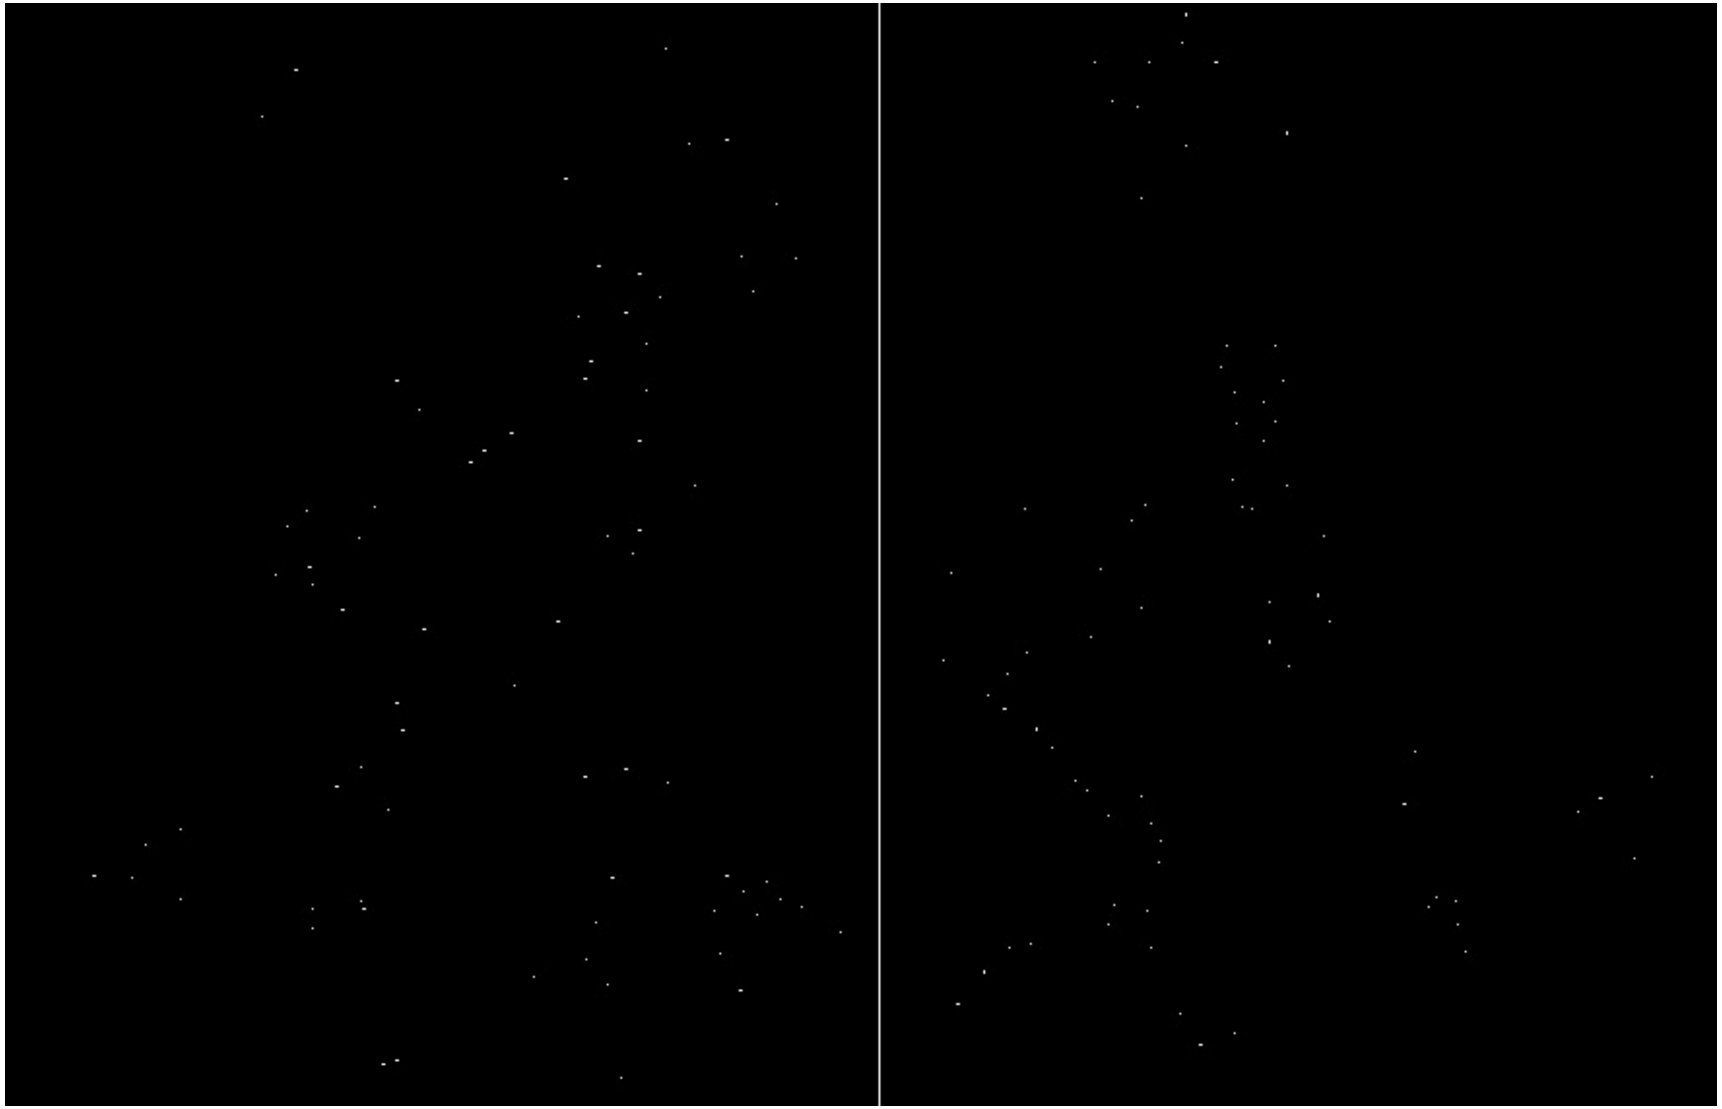

Harris Detector: Steps

Harris Corners – Why so complicated?

- Can't we just check for regions with lots of gradients in the x and y directions (or any specific)?
  - No! A diagonal line or alike would satisfy that criteria

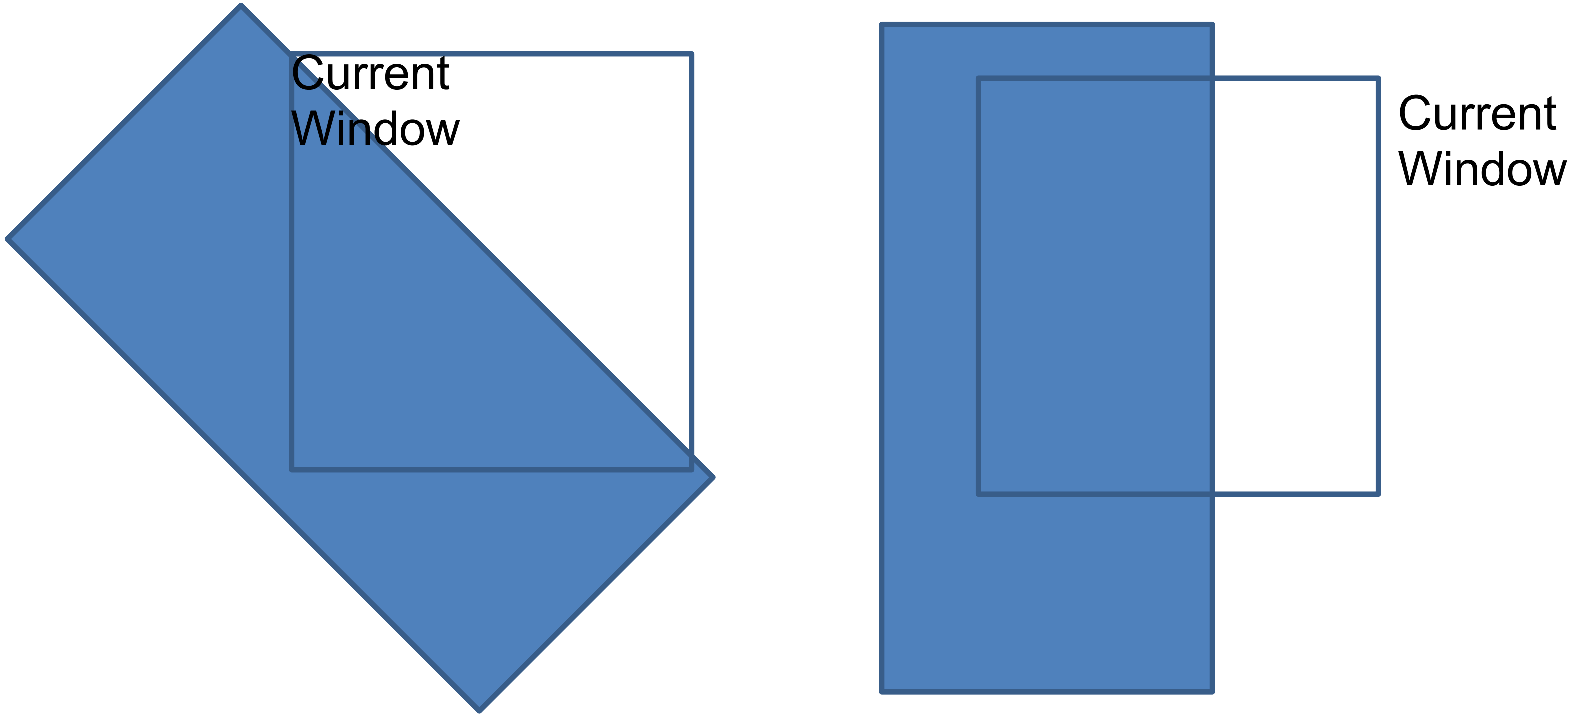

Invariance and covariance

Are locations *invariant* to photometric transformations and *covariant* to geometric transformations?

- **Invariance:** image is transformed and corner locations do not change
- **Covariance:** if we have two transformed versions of the same image, features should be detected in corresponding locations

Affine intensity change

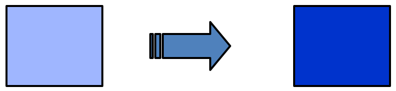

$$I \rightarrow a\,I + b$$

•Only derivatives are used =>
invariance to intensity shift $I \rightarrow I + b$

- Intensity scaling: $I \rightarrow a\,I$

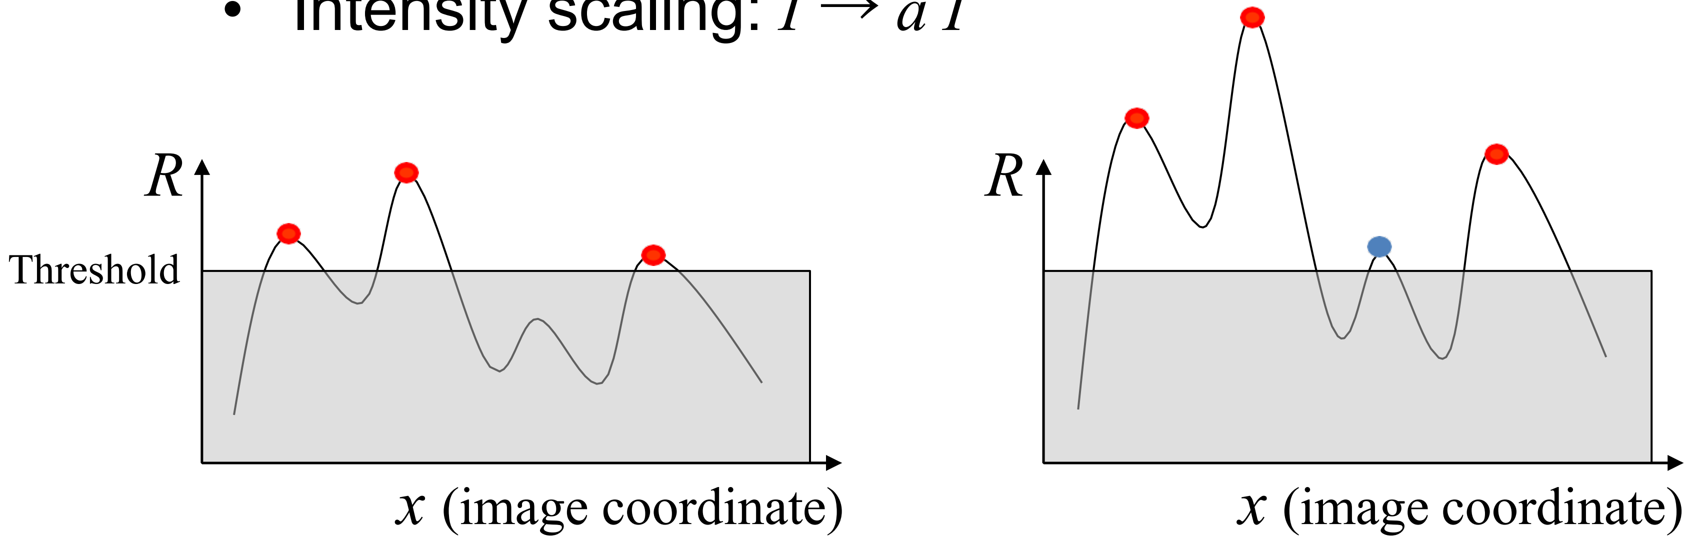

Partially invariant to affine intensity change

Image translation

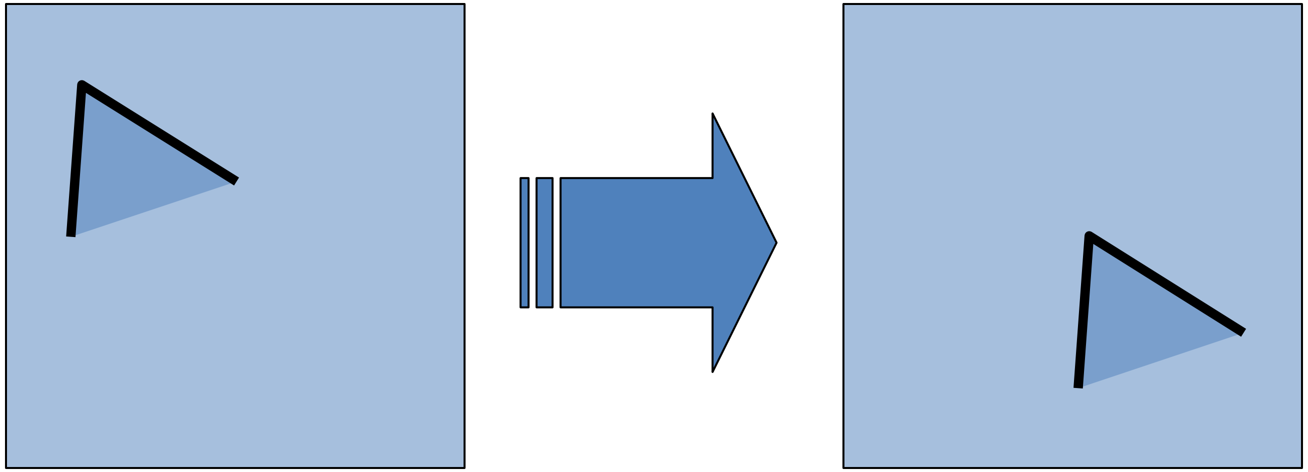

- Derivatives and window function are shift-invariant.

Corner location is covariant w.r.t. translation

James Hays

Image rotation

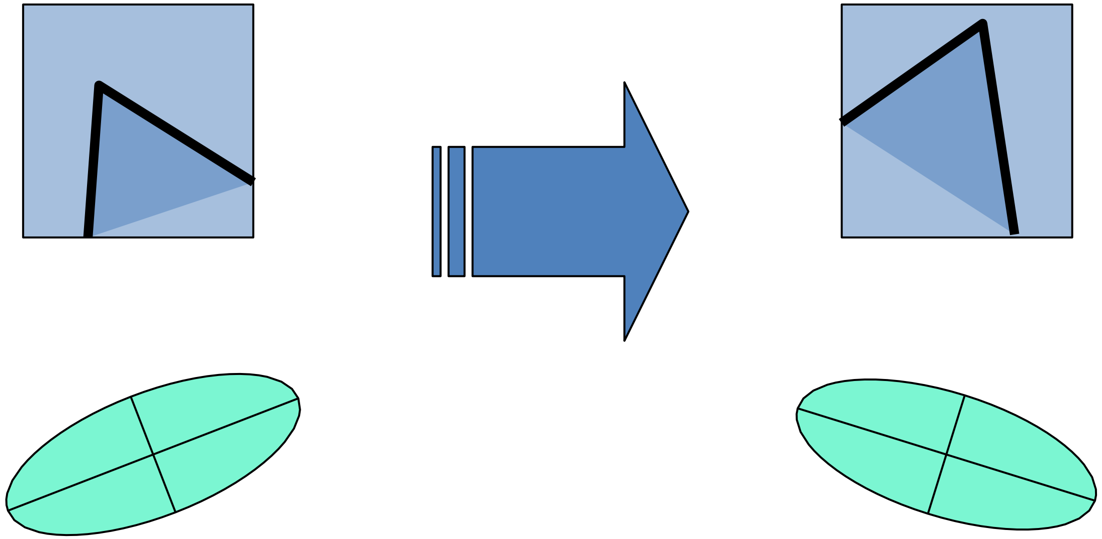

Second moment ellipse rotates but its shape (i.e., eigenvalues) remains the same.

Corner location is covariant w.r.t. rotation

James Hays

Scaling

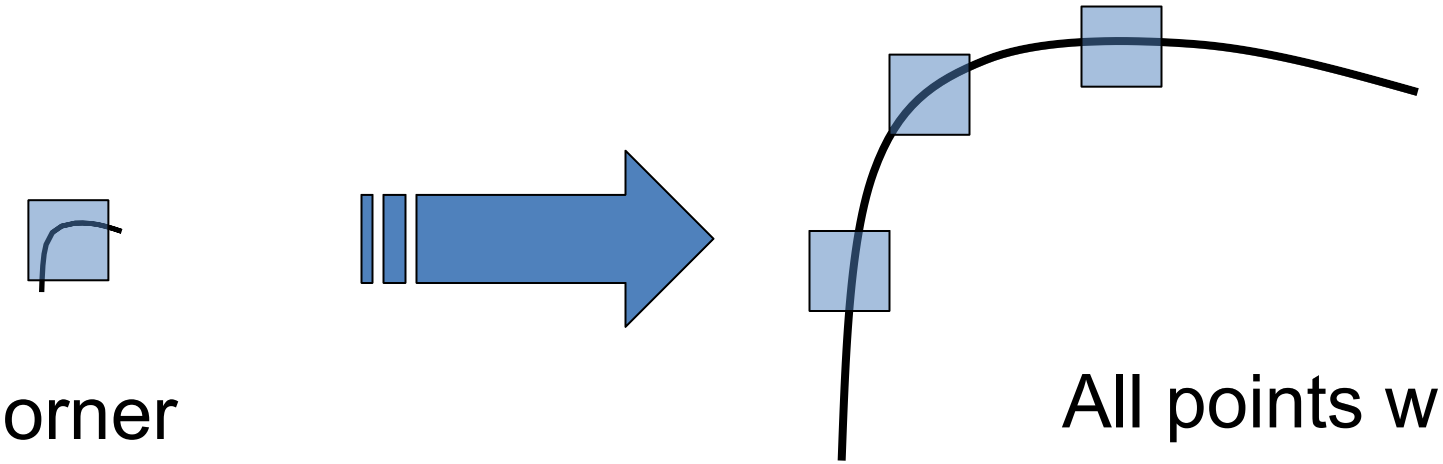

Corner

All points will be incorrectly classified as edges

Corner location is not covariant to scaling!

James Hays

Automatic Scale Selection

$$f(I_{i_1\ldots i_m}(x,a)) = f(I_{i_1\ldots i_m}(x',a'))$$

How to find patch sizes at which *f* response is equal?

What is a good *f* ?

K. Grauman, B. Leibe

Automatic Scale Selection

- Function responses for increasing scale (scale signature)

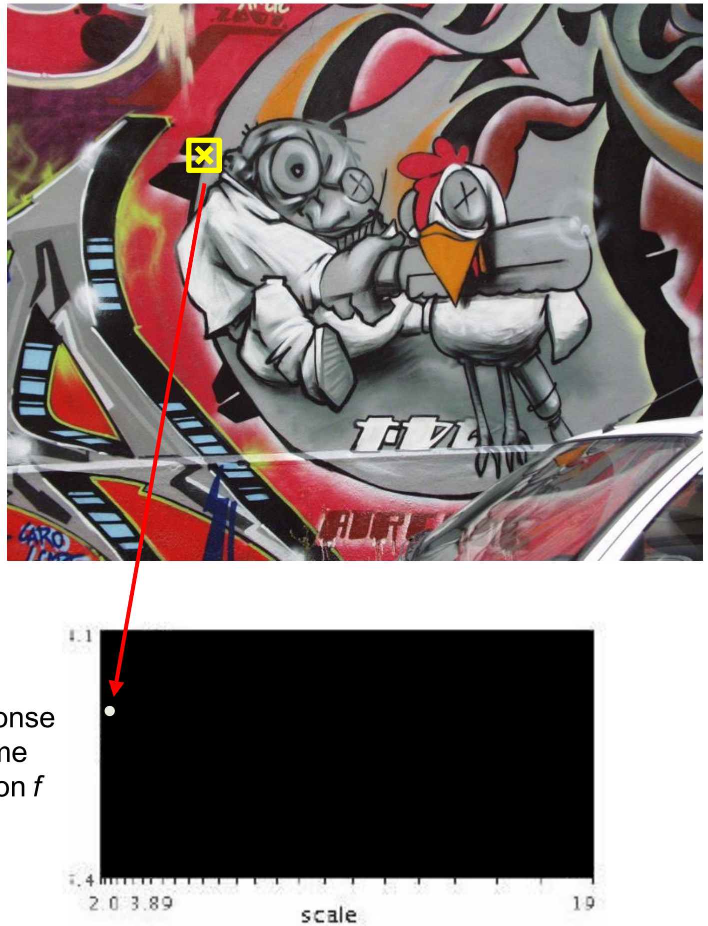

Response of some function *f*

$f(I_{i_1 \ldots i_m}(x, \quad))$

$f(I_{i_1 \ldots i_m}(x', \quad'))$

K. Grauman, B. Leibe

Automatic Scale Selection

- Function responses for increasing scale (scale signature)

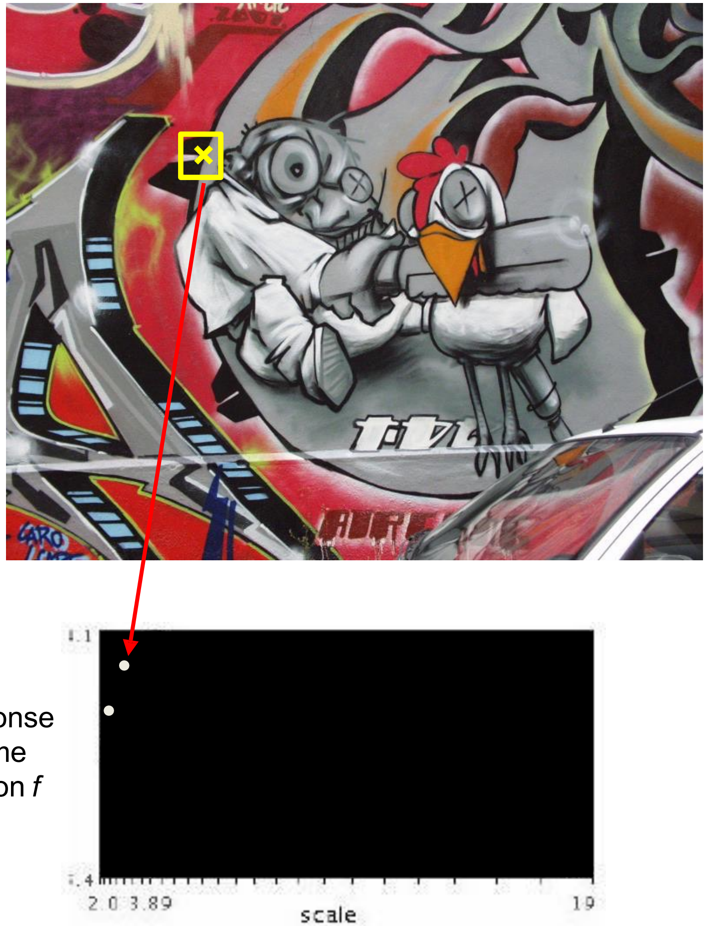

Response
of some
function *f*

$f(I_{i_1 \ldots i_m}(x, \ ))$

$f(I_{i_1 \ldots i_m}(x', \ '))$

K. Grauman, B. Leibe

Automatic Scale Selection

- Function responses for increasing scale (scale signature)

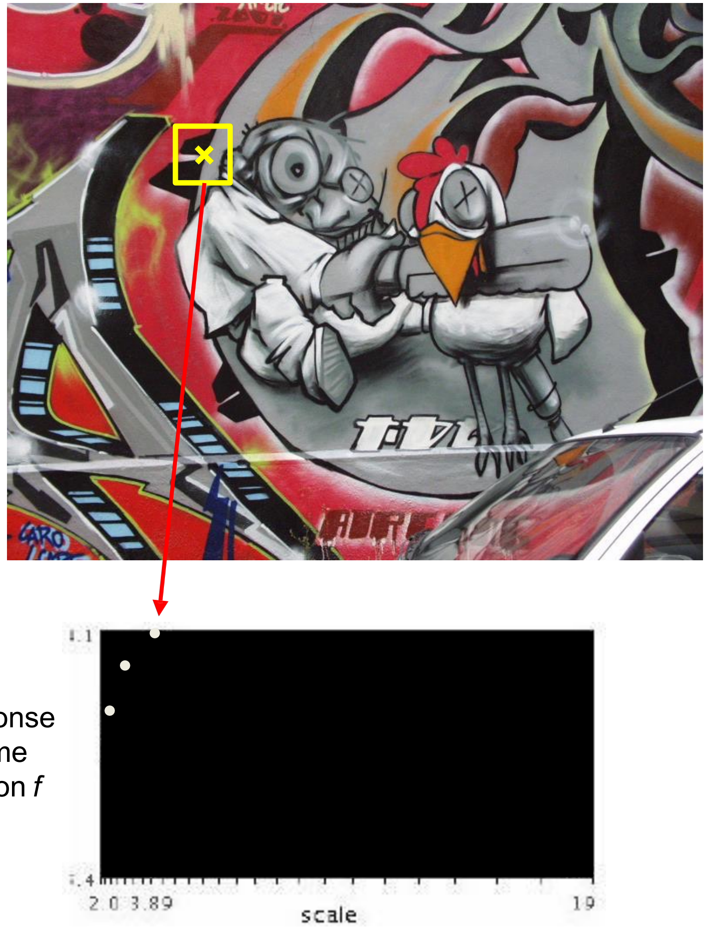

Response of some function *f*

$f(I_{i_1 \ldots i_m}(x, \ ))$

$f(I_{i_1 \ldots i_m}(x', \ '))$

K. Grauman, B. Leibe

Automatic Scale Selection

- Function responses for increasing scale (scale signature)

Response of some function *f*

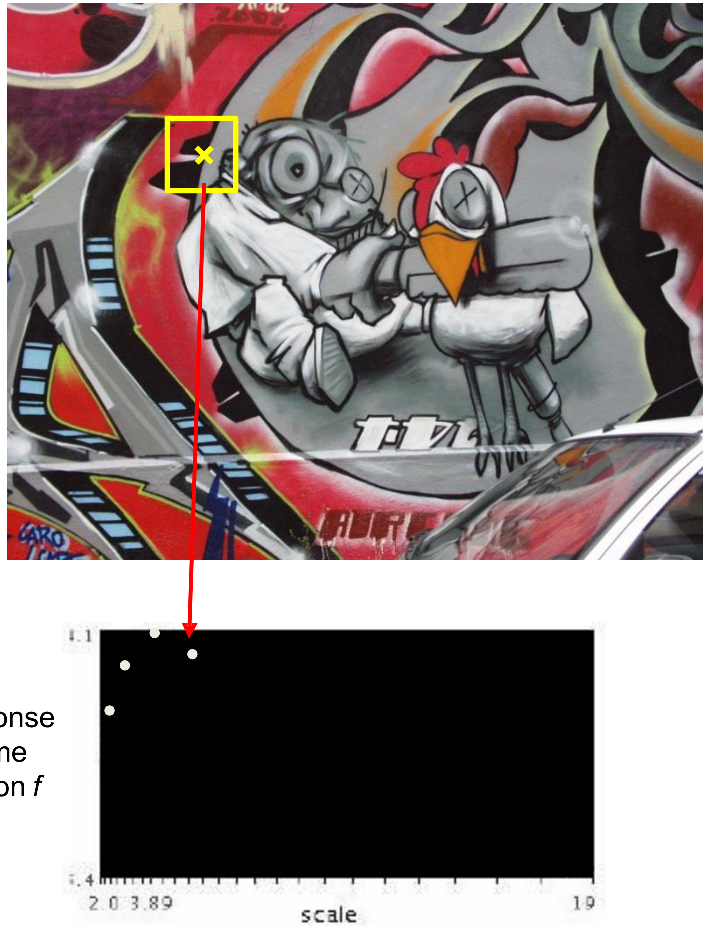

$f(I_{i_1 \ldots i_m}(x, \ ))$

$f(I_{i_1 \ldots i_m}(x', \ '))$

K. Grauman, B. Leibe

Automatic Scale Selection

- Function responses for increasing scale (scale signature)

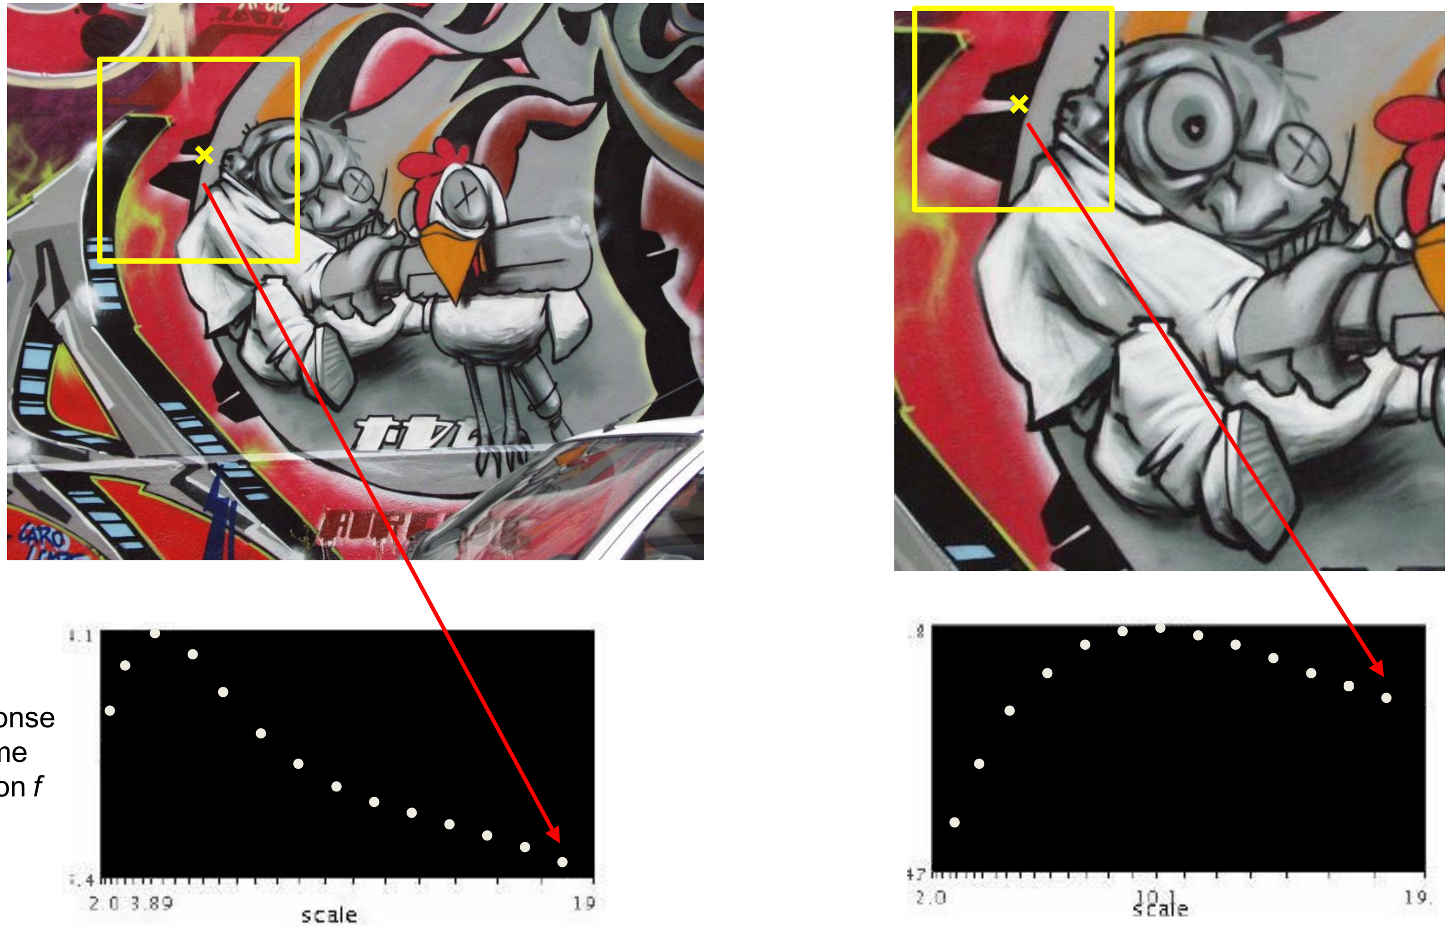

Response of some function *f*

$f(I_{i_1\ldots i_m}(x,\ ))$

$f(I_{i_1\ldots i_m}(x',\ '))$

K. Grauman, B. Leibe

Automatic Scale Selection

- Function responses for increasing scale (scale signature)

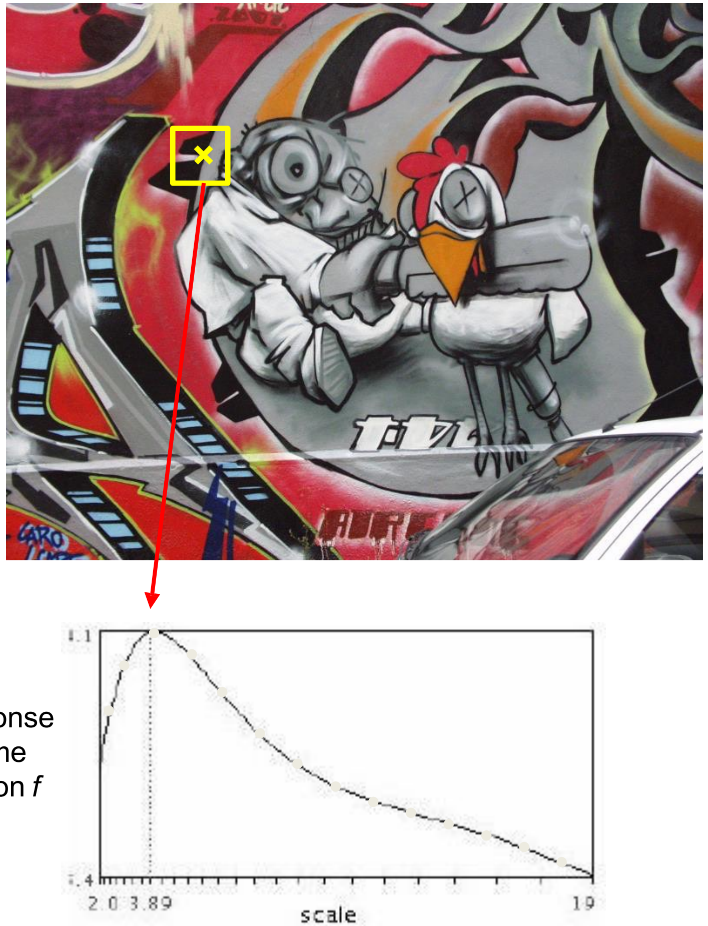

Response
of some
function *f*

$$f(I_{i_1 \ldots i_m}(x, \ ))$$

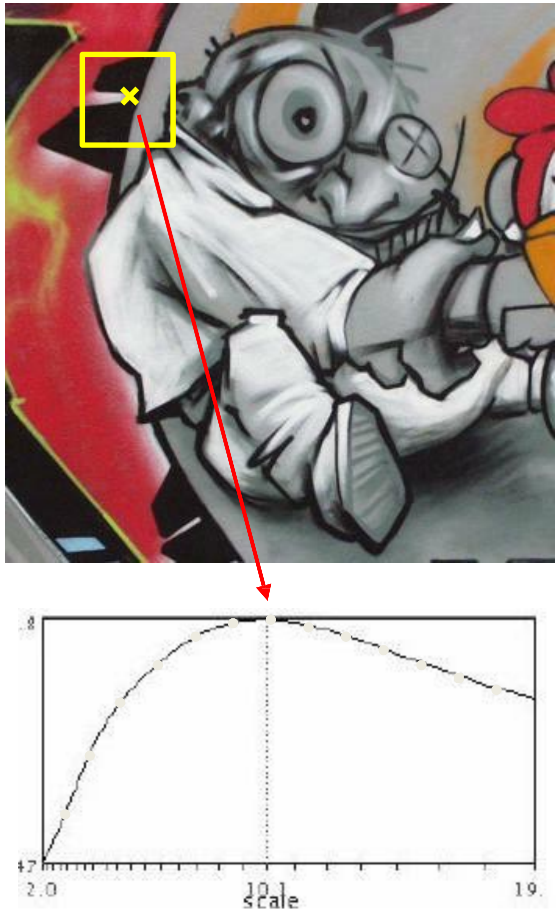

$$f(I_{i_1 \ldots i_m}(x', \ '))$$

What Is A Useful Signature Function *f* ?

- "Blob" detector is common for corners
  - Laplacian (2nd derivative) of Gaussian (LoG)

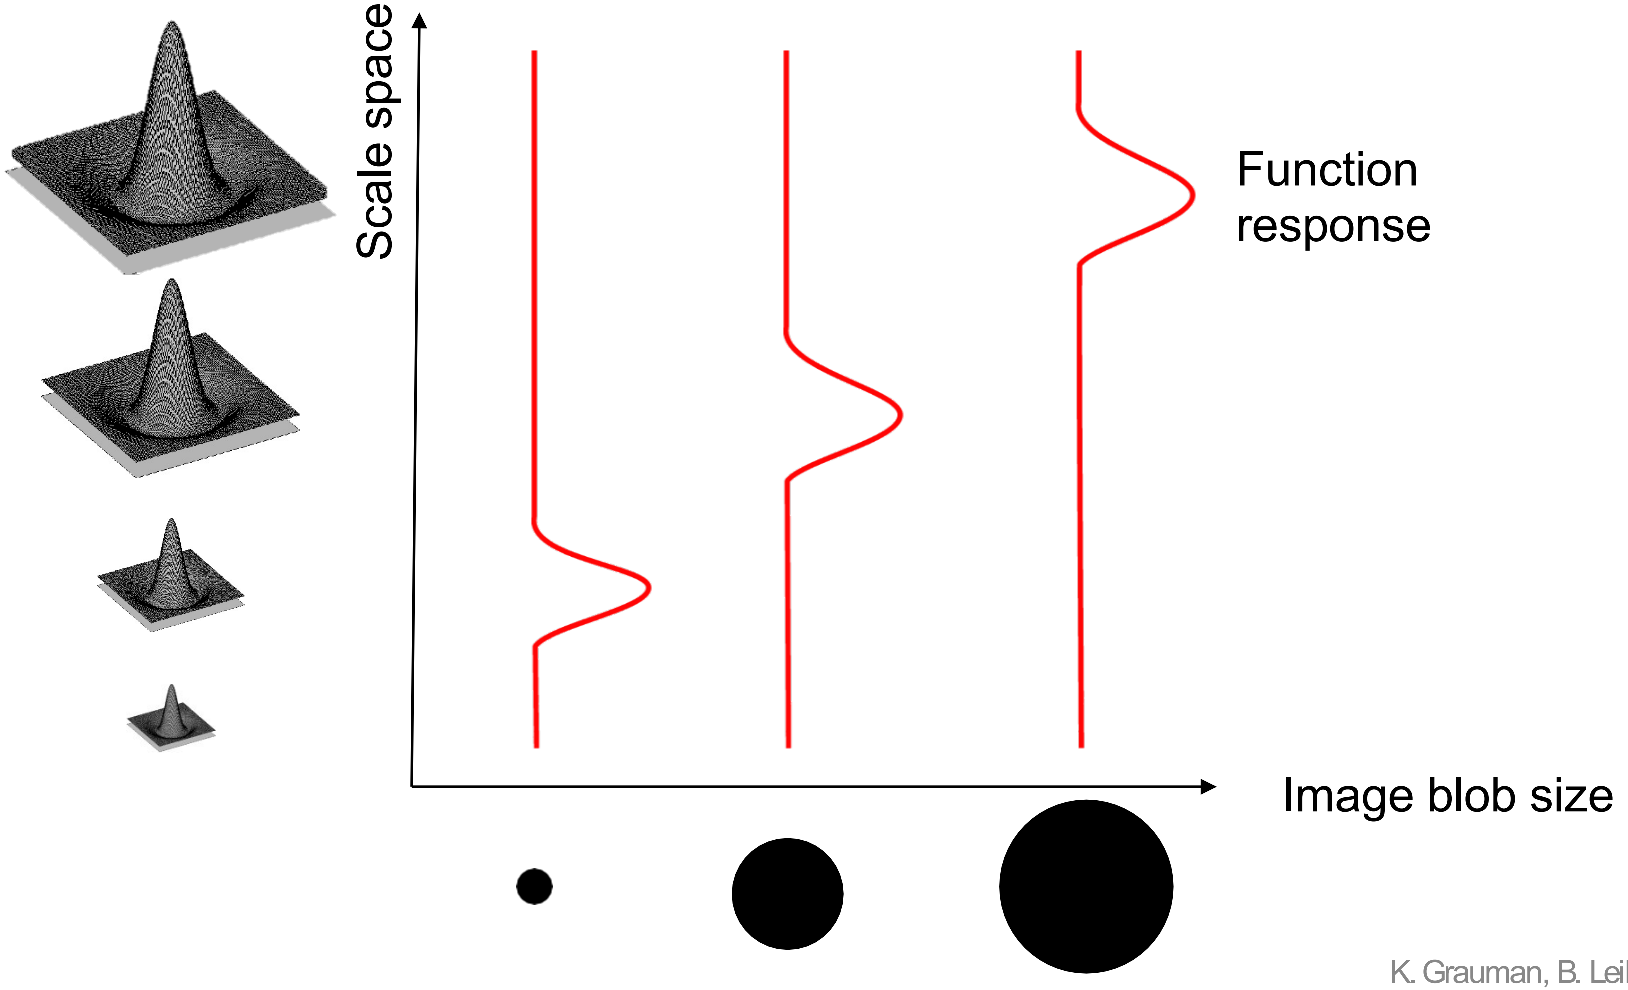

K. Grauman, B. Leibe

Harris-Laplace detector [Mikolajczyk '01]

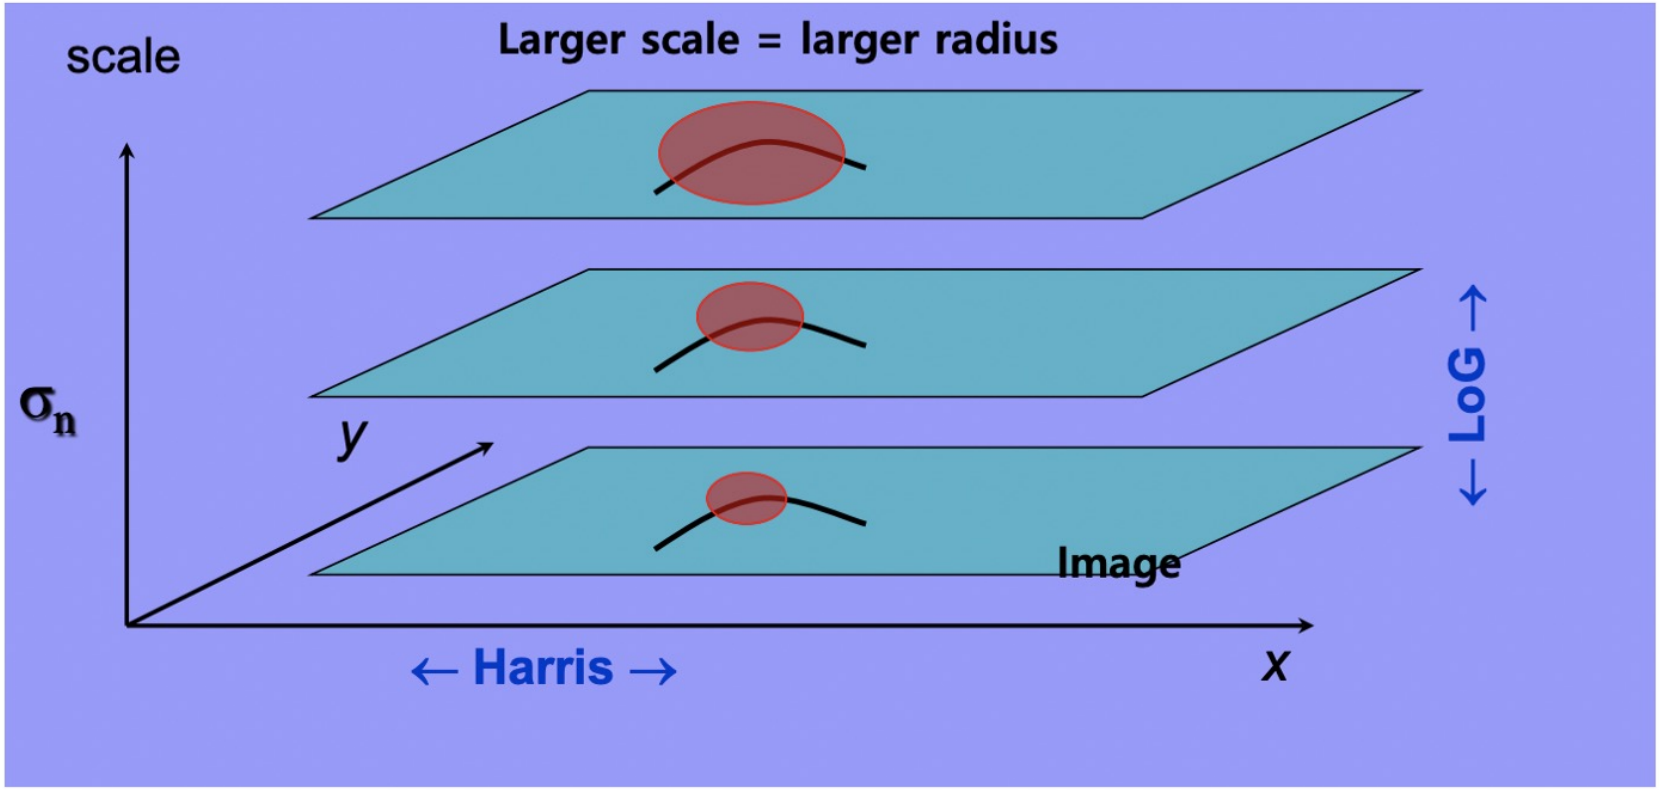

- **Step 1.** Build the Laplacian Pyramid of one image

Harris-Laplace detector [Mikolajczyk ‘01]

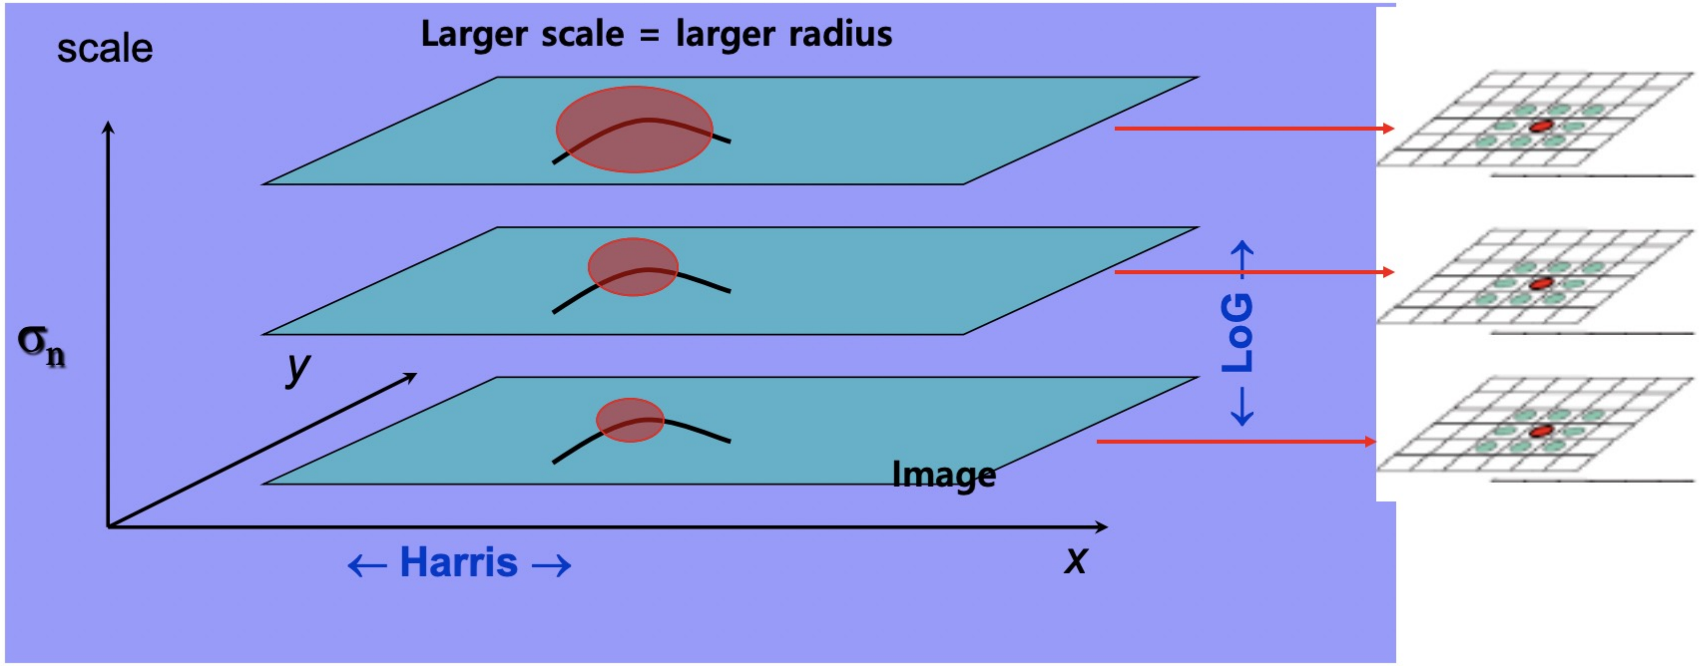

- **Step 1.** Build the Laplacian Pyramid of one image
- **Step 2**. Run the *Harris* detector to compute interest points *at each scale*

Harris-Laplace detector [Mikolajczyk ‘01]

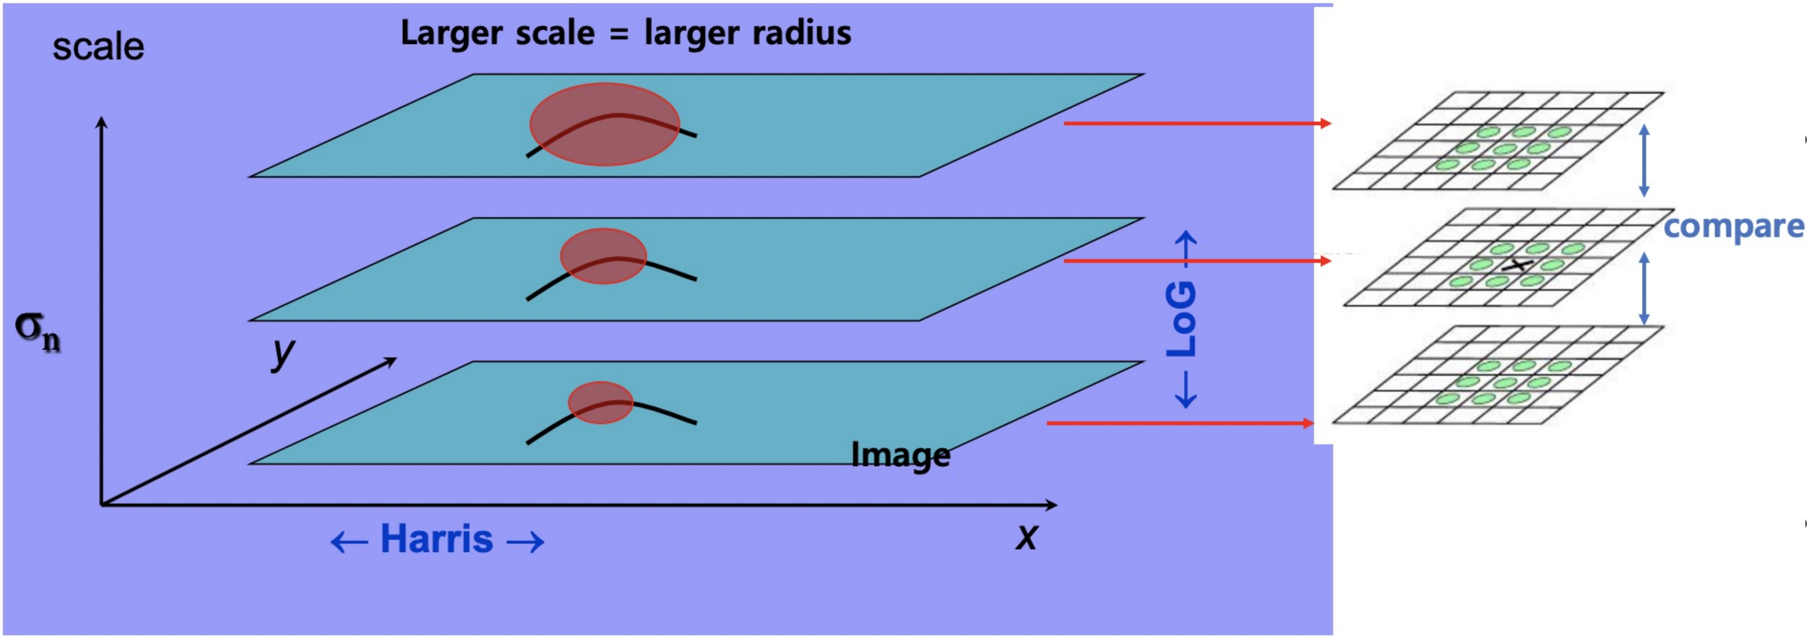

- **Step 1.** Build the Laplacian Pyramid of one image
- **Step 2**. Run the *Harris* detector to compute interest points *at each scale*
- **Step 3**. Non-maximal suppression, not only at each scale, **but also at adjacent scales**

Harris-Laplace detector [Mikolajczyk ‘01]

- **A scale-invariant detector!**
  - *Automatically search for the right scale to detect corners, by “multi-scaling then max-pooling”*

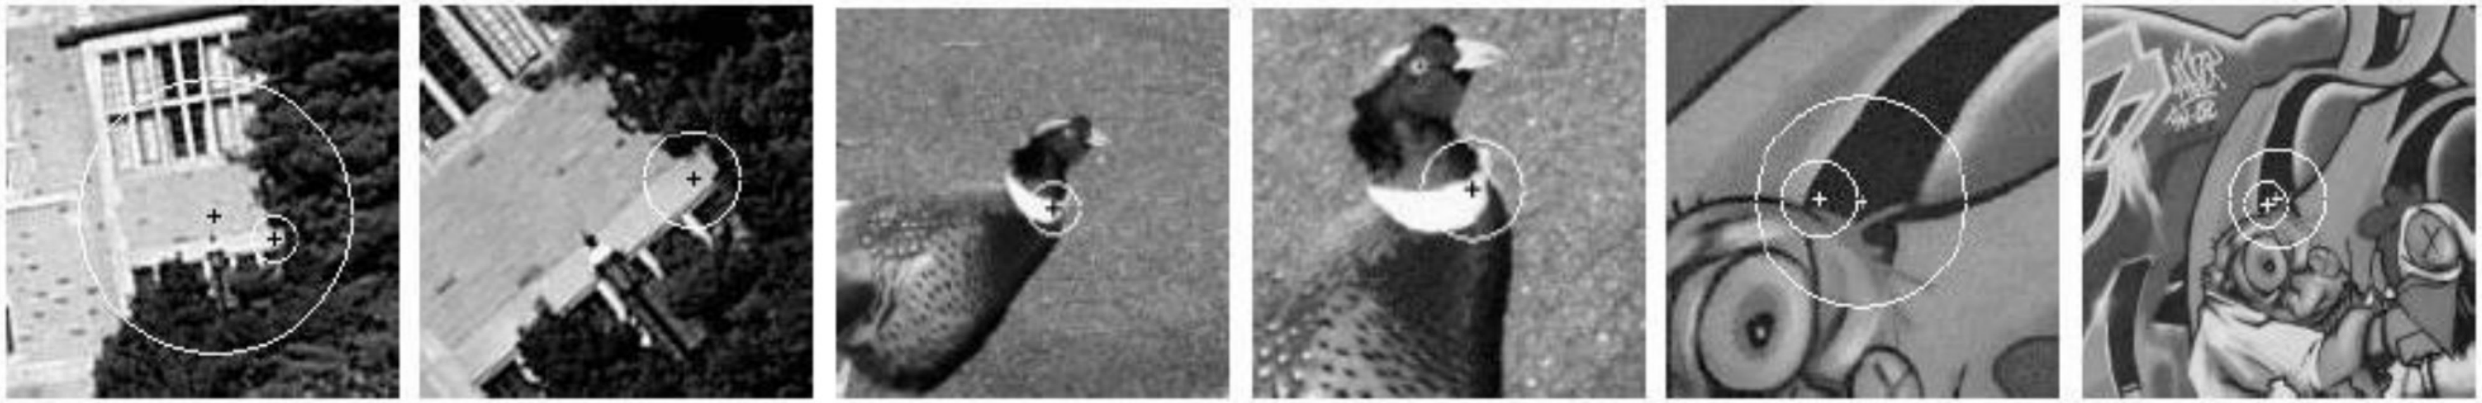

Harris-Laplace points

A Longer List of Local Keypoint Detectors…

Table 7.1 Overview of feature detectors.

Feature Detector	Corner	Blob	Region	Rotation invariant	Scale invariant	Affine invariant	Repeatability	Localization accuracy	Robustness	Efficiency
Harris	√			√			+++	+++	+++	++
Hessian		√		√			++	++	++	+
SUSAN	√			√			++	++	++	+++
Harris-Laplace	√	(√)		√	√		+++	+++	++	+
Hessian-Laplace	(√)	√		√	√		+++	+++	+++	+
DoG	(√)	√		√	√		++	++	++	++
SURF	(√)	√		√	√		++	++	++	+++
Harris-Affine	√	(√)		√	√	√	+++	+++	++	++
Hessian-Affine	(√)	√		√	√	√	+++	+++	+++	++
Salient Regions	(√)	√		√	√	(√)	+	+	++	+
Edge-based	√			√	√	√	+++	+++	+	+
MSER			√	√	√	√	+++	+++	++	+++
Intensity-based			√	√	√	√	++	++	++	++
Superpixels			√	√	(√)	(√)	+	+	+	+

- It is always of interest to collect more points with more detectors, for more possible matches
- Need consider location preciseness, variation robustness, and flexibility in region shapes
- Best choice often application dependent
  - Harris/Harris-Laplace work well for many natural image categories
  - MSER works well for buildings and printed things
  - Although no “silver bullet”, all detectors/descriptors shown here work well in general

Tuytelaars Mikolajczyk 2008

Local features: main components

1) Detection:

Find a set of distinctive key points.

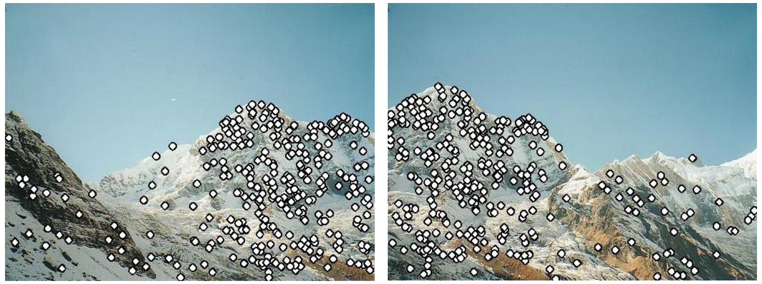

2) Description:

Extract feature descriptor around each interest point as vector.

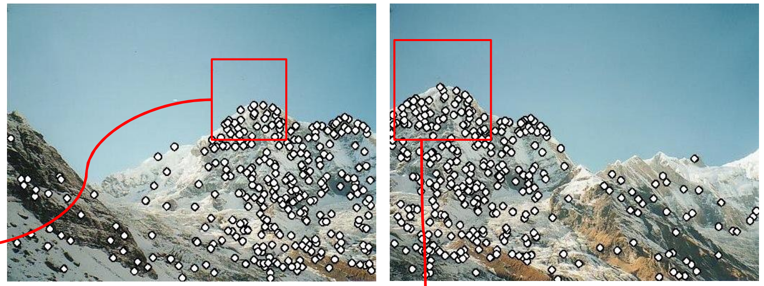

x_1 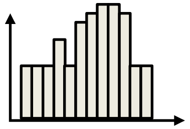

$$\mathbf{x}_1 = [x_1^{(1)}, \ldots, x_d^{(1)}]$$

x_2 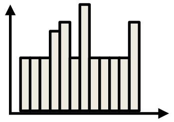

$$\mathbf{x}_2 = [x_1^{(2)}, \ldots, x_d^{(2)}]$$

3) Matching:

Compute distance between feature vectors to find correspondence.

$$d(\mathrm{x}_1, \mathrm{x}_2) < T$$

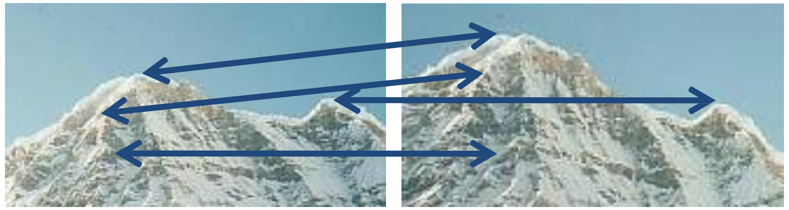

Image representations

- Templates
  - Intensity, gradients, etc.

- Histograms
  - Color, texture, SIFT descriptors, etc.

James Hays

Image patch

Just use the pixel values of the patch

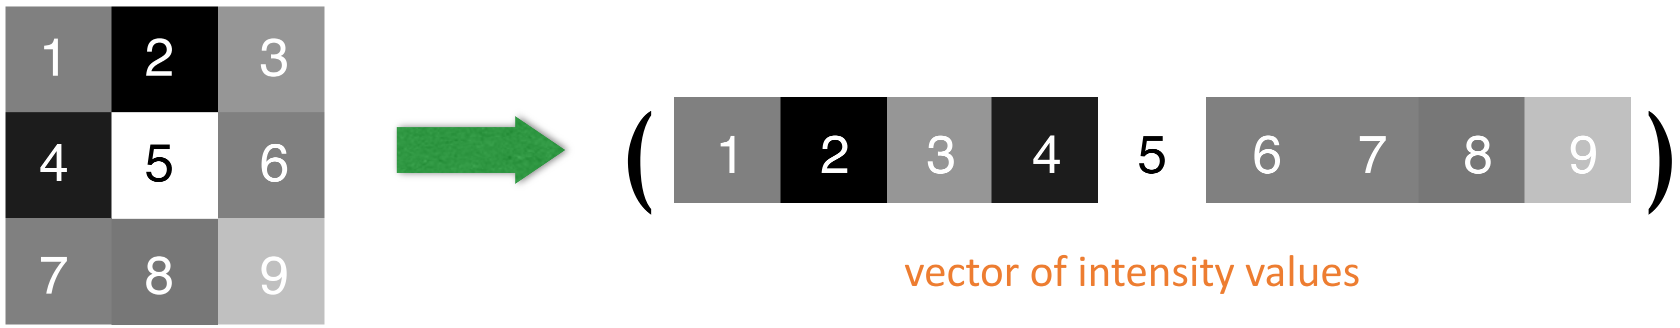

Perfectly fine if geometry and appearance is unchanged (a.k.a. template matching)

What are the problems?

How can you be less sensitive to absolute intensity values?

Image gradients

Use pixel differences

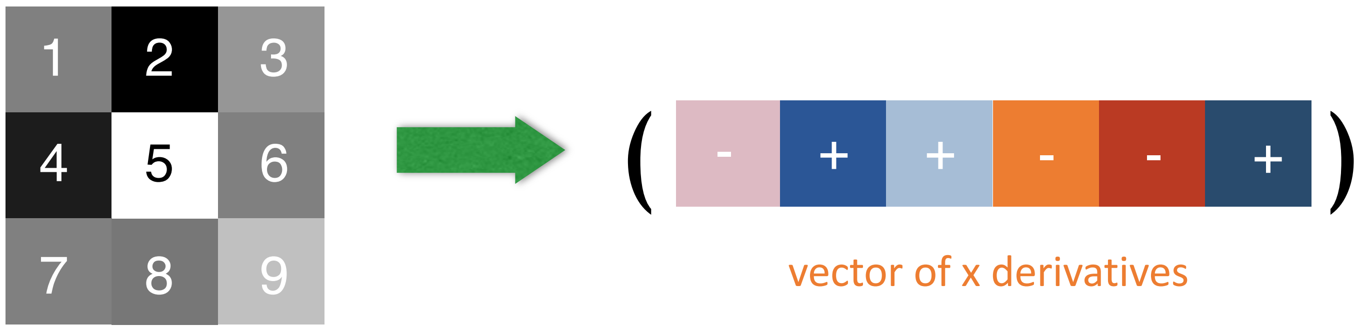

vector of x derivatives

Feature is invariant to absolute intensity values

What are the problems?

How can you be less sensitive to deformations?

Image Representations: Histograms

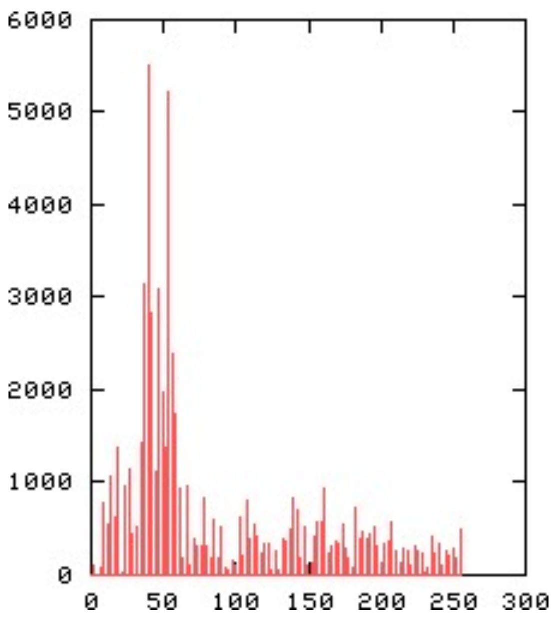

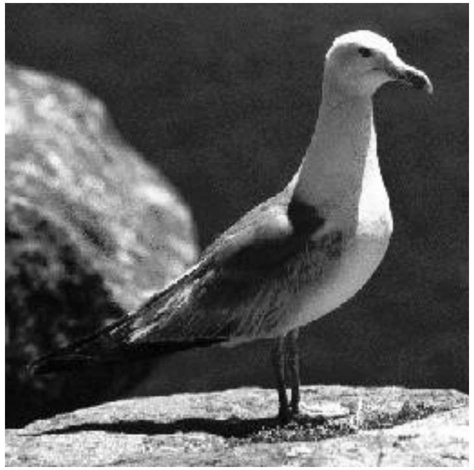

Motivation: We want some sensitivity to spatial layout, but not too much, so blocks of histograms give us that

Global histogram to represent distribution of features

– How 'well exposed' a photo is

What about a local histogram per detected point?

Images from Dave Kauchak

Image Representations: Histograms

Histogram: Probability or count of data in each bin

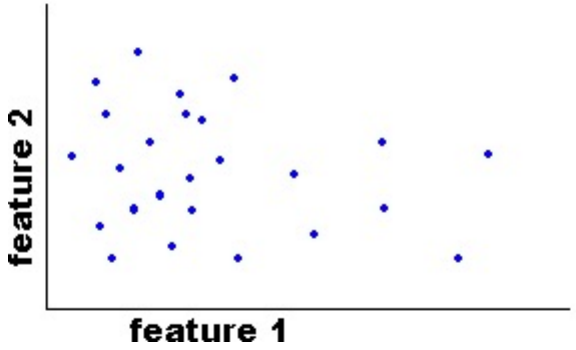

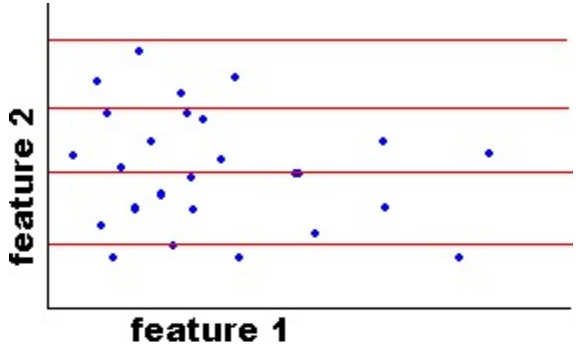

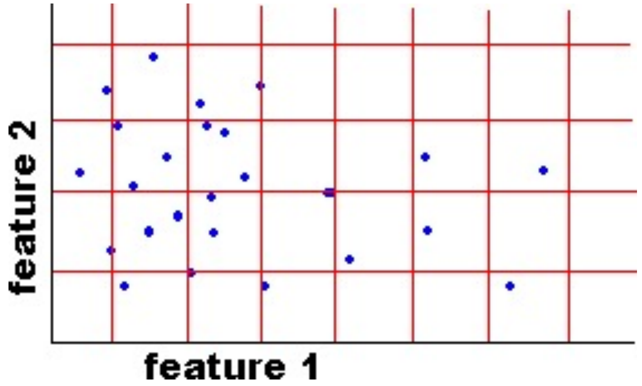

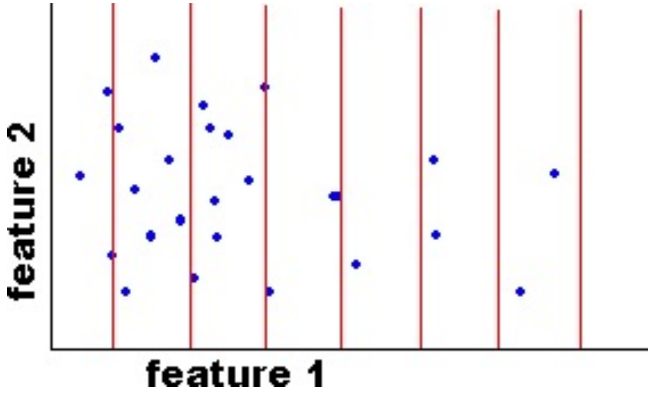

- Joint histogram
  - Requires lots of data
  - Loss of resolution to avoid empty bins

Marginal histogram

- Requires independent features
- More data/bin than joint histogram
- **The definition of feature can vary…**

Images from Dave Kauchak

Image Representations: Histograms

Clustering

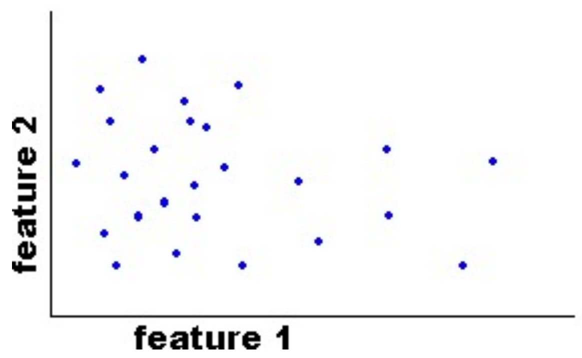

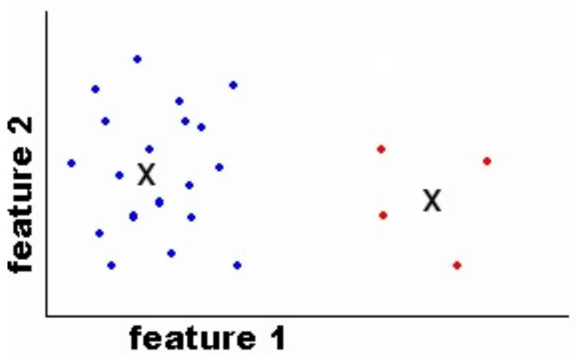

Use the same cluster centers for all images
… so two different images' histograms can be "compared" (using what metric?)

Images from Dave Kauchak

Computing histogram distance

$$\text{histint}(h_i, h_j) = 1 - \sum_{m=1}^{K} \min\left(h_i(m), h_j(m)\right)$$

Histogram intersection (assuming normalized histograms)

$$\chi^2(h_i, h_j) = \frac{1}{2} \sum_{m=1}^{K} \frac{[h_i(m) - h_j(m)]^2}{h_i(m) + h_j(m)}$$

Chi-squared Histogram matching distance

Cars found by color histogram matching using chi-squared

James Hays

Histograms: Implementation issues

- Quantization
  - Grids: fast but applicable only with few dimensions
  - Clustering: slower but can quantize data in higher dimensions

Few Bins
Need less data
Coarser representation

Many Bins
Need more data
Finer representation

- Matching
  - Histogram intersection or Euclidean may be faster
  - Chi-squared often works better
  - Earth mover's distance is good for when nearby bins represent similar values

Intensity/Color histogram

Count the colors in the image using a histogram

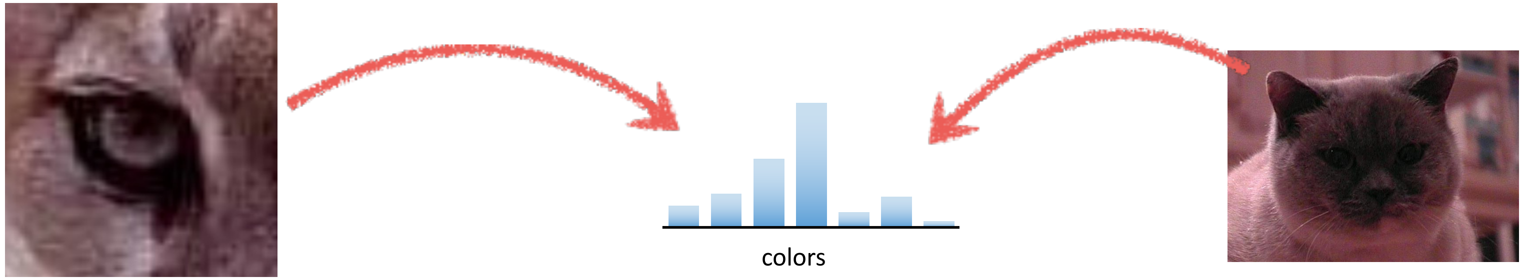

Invariant to changes in scale and rotation

What are the problems?

How can you be more sensitive to spatial layout?

Spatial histograms

Compute histograms over spatial 'cells'

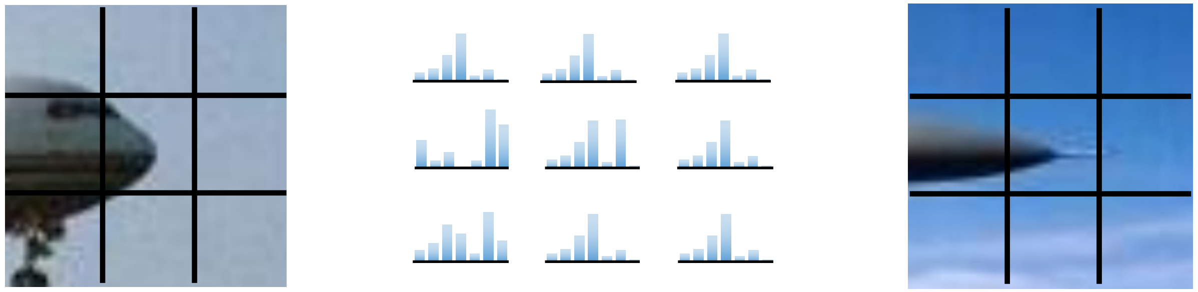

Retains rough spatial layout

Some invariance to deformations

What are the problems?

How can you be completely invariant to rotation?

For what things might we compute histograms?

- *From Color to Texture/Keypoints*
- Histograms of oriented gradients (**HOG**)
- Scale Invariant Feature Transform (**SIFT**)
  - Extremely popular (63k citations in 2021)

IMHO, one of the most elegant designs ever in CV

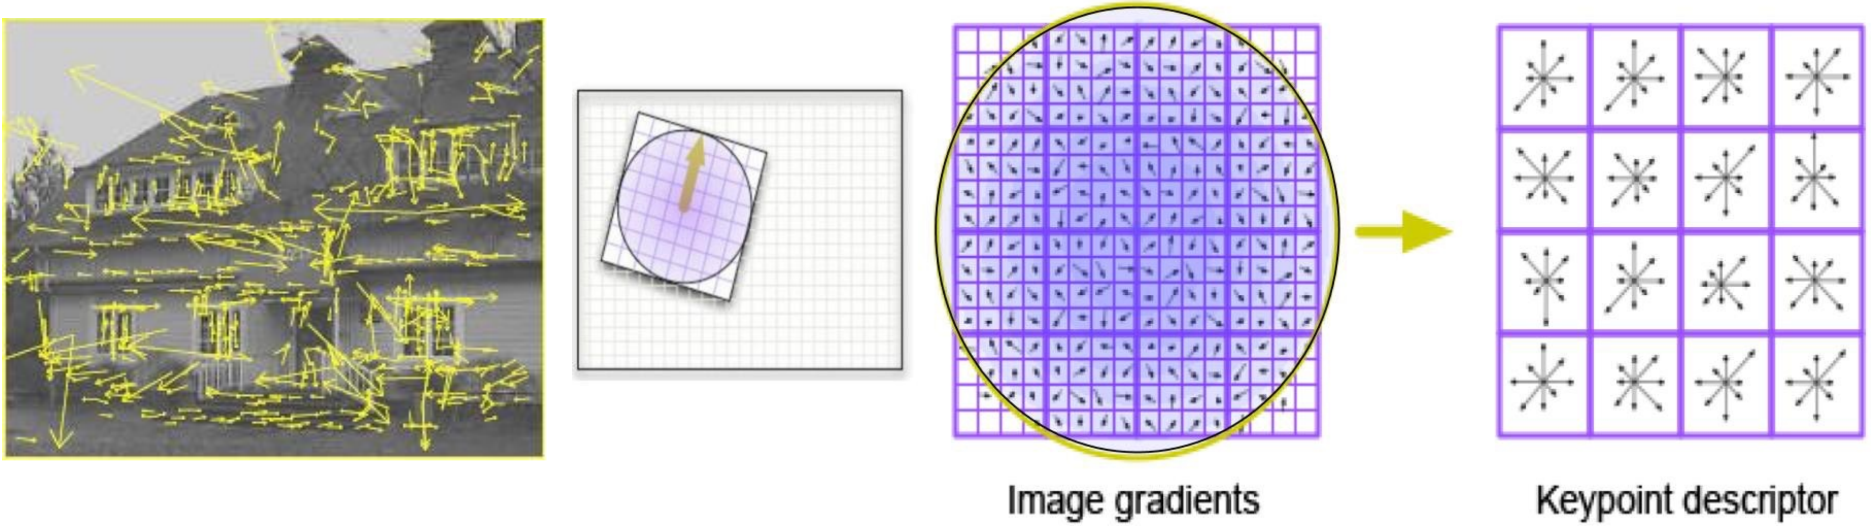

SIFT – Lowe IJCV 2004

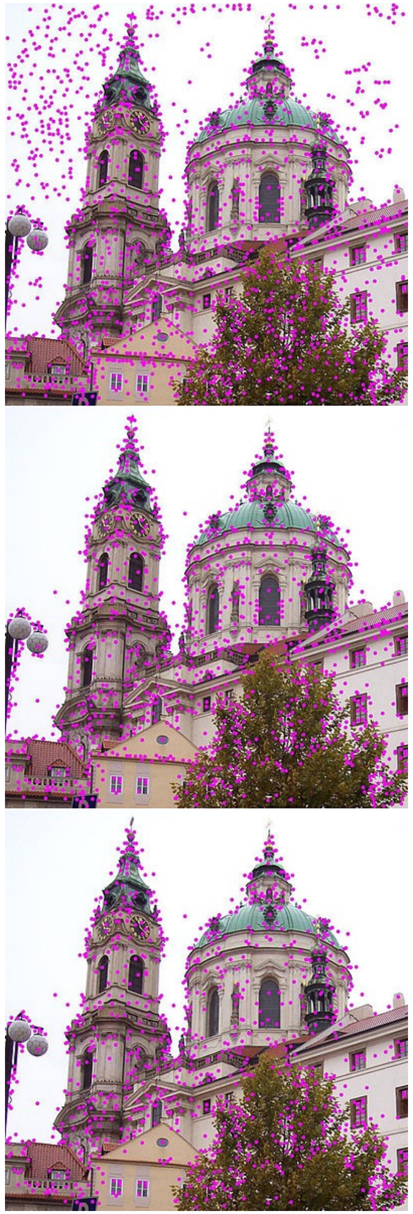

SIFT (Scale Invariant Feature Transform)

SIFT describes both a **detector** and **descriptor**

1. Multi-scale extrema detection
2. Keypoint localization
3. Orientation assignment
4. Keypoint descriptor

1. Multi-scale extrema detection

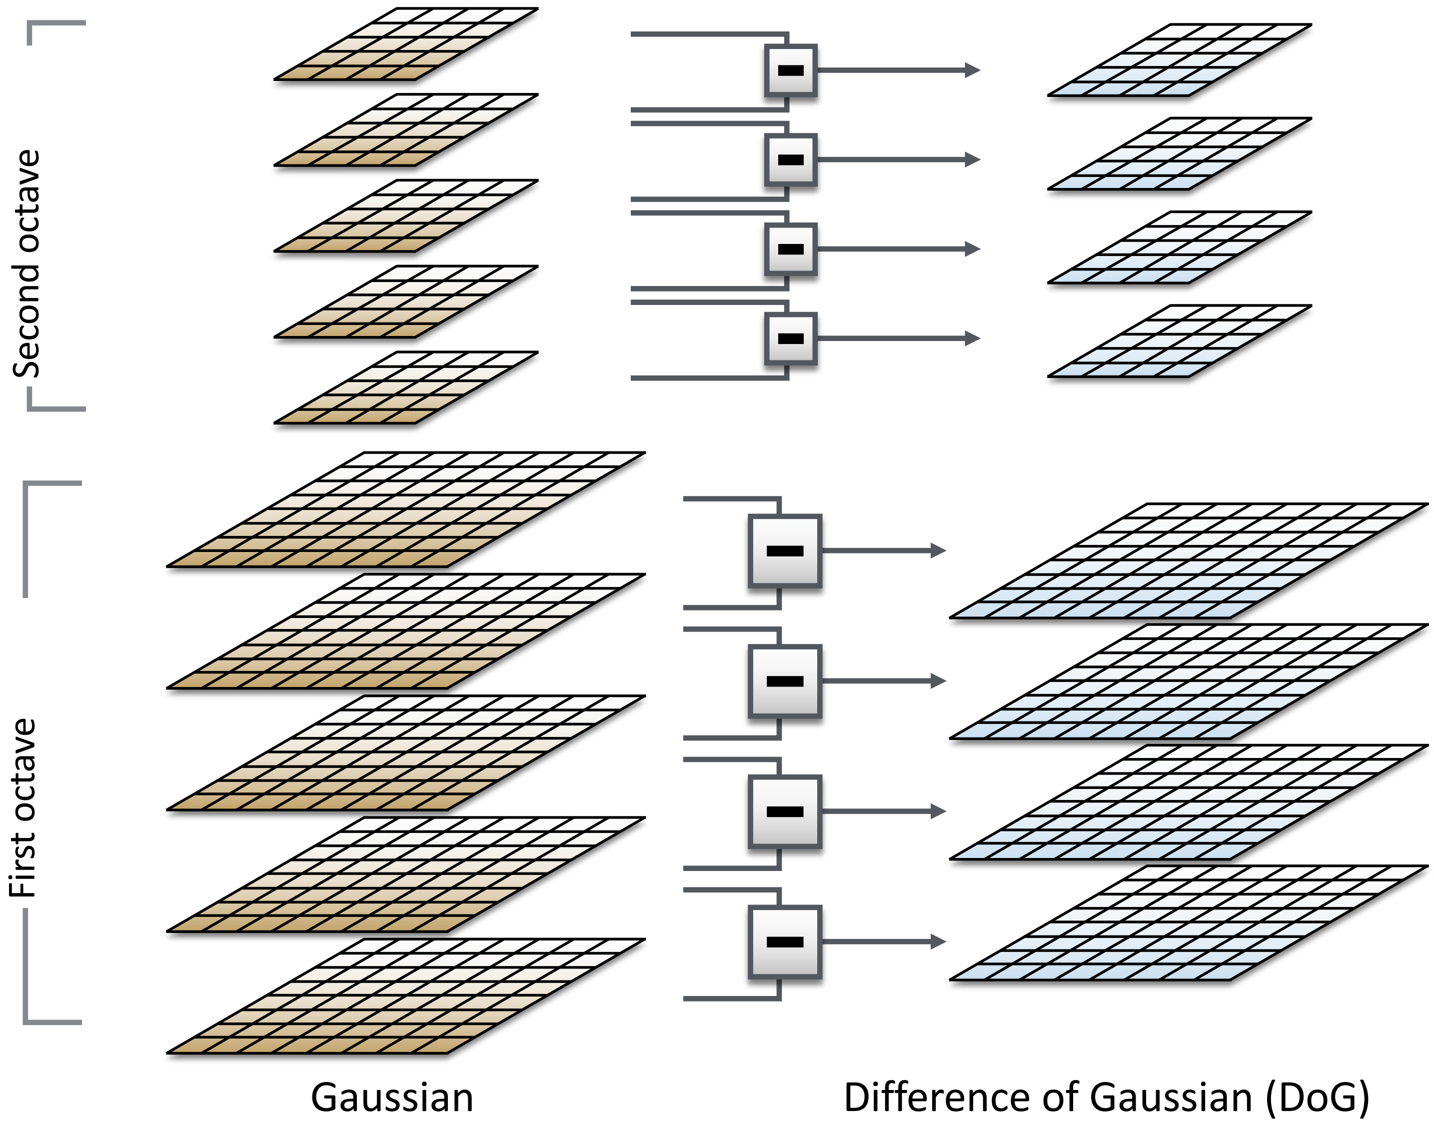

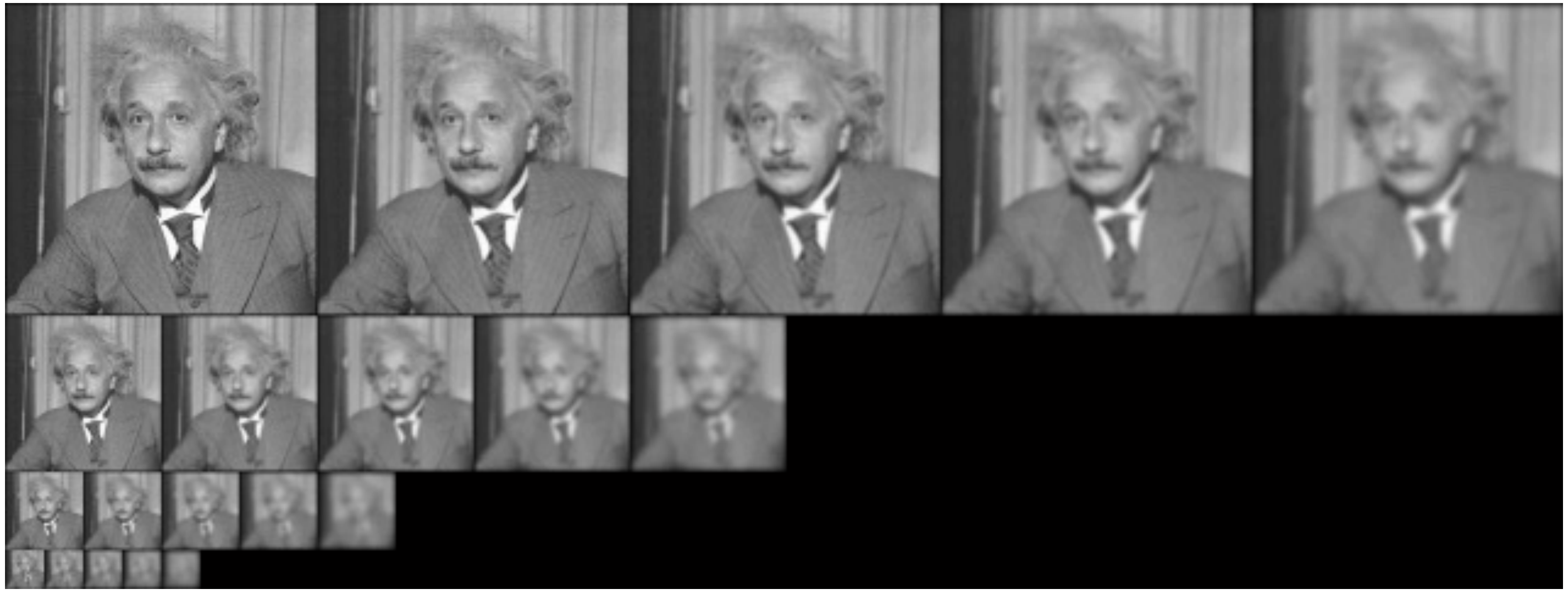

Gaussian

Laplacian

Scale-space extrema

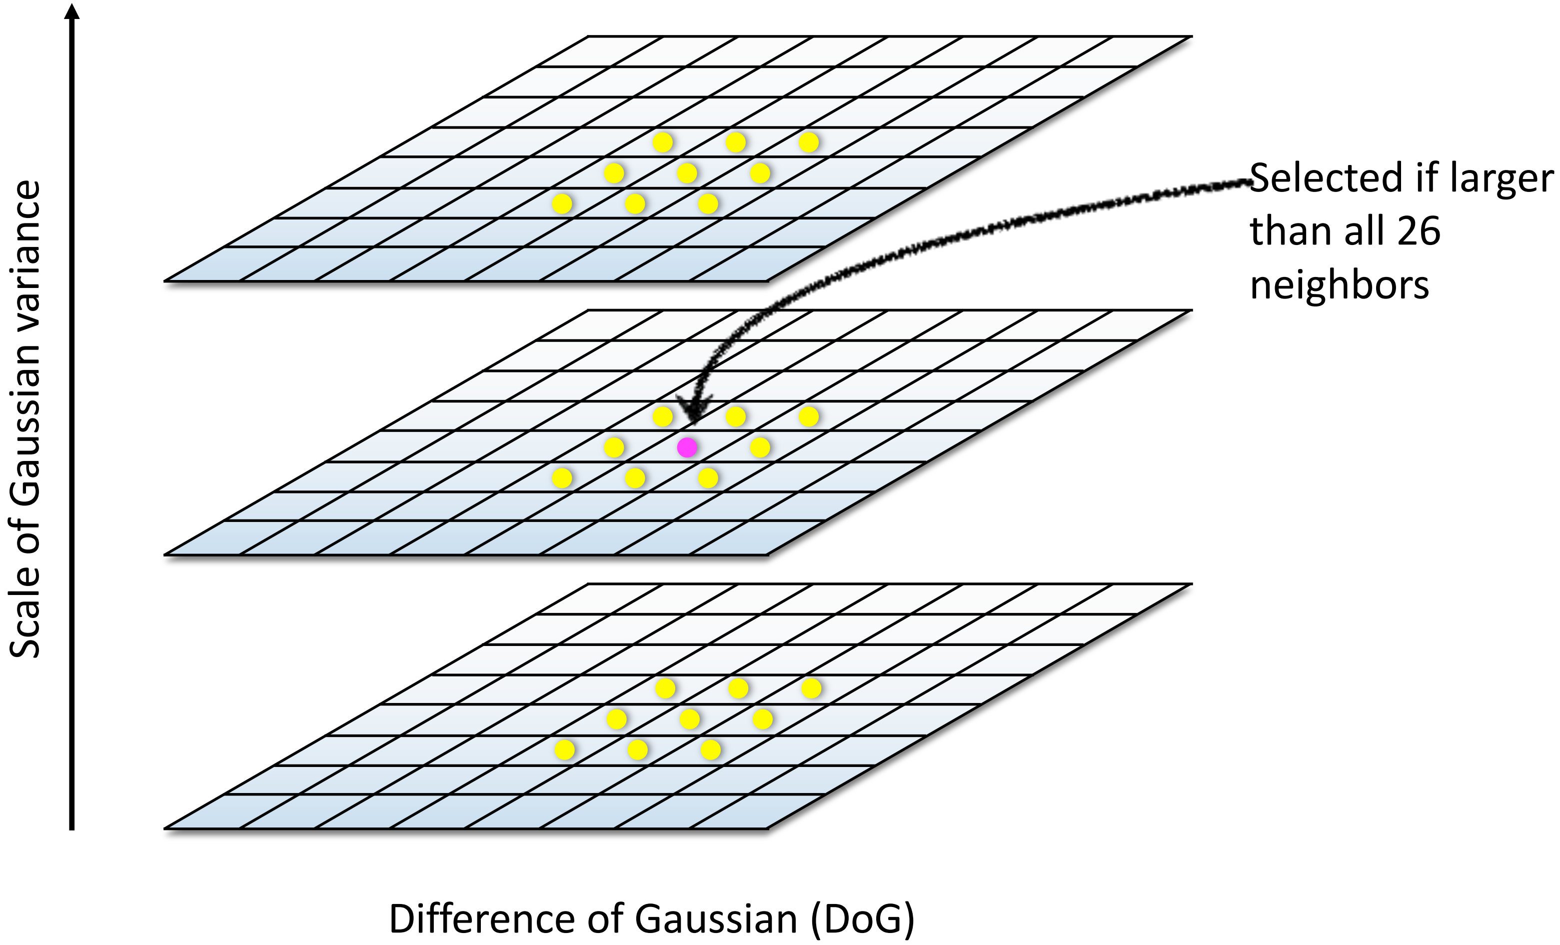

2. Keypoint localization (why this step?)

- Reject flats:
  - $|D(\hat{\mathbf{x}})|$ < 0.03
- Reject edges:

$$\mathbf{H} = \begin{bmatrix} D_{xx} & D_{xy} \\ D_{xy} & D_{yy} \end{bmatrix}$$

Let α be the eigenvalue with larger magnitude and β the smaller.

$$\mathrm{Tr}(\mathbf{H}) = D_{xx} + D_{yy} = \alpha + \beta,$$

$$\mathrm{Det}(\mathbf{H}) = D_{xx}D_{yy} - (D_{xy})^2 = \alpha\beta.$$

Let r = α/β.
So α = rβ

$$\frac{\mathrm{Tr}(\mathbf{H})^2}{\mathrm{Det}(\mathbf{H})} = \frac{(\alpha+\beta)^2}{\alpha\beta} = \frac{(r\beta+\beta)^2}{r\beta^2} = \frac{(r+1)^2}{r},$$

$(r+1)^2/r$ is at a min when the 2 eigenvalues are equal.

  - r < 10

3. Orientation assignment

For a keypoint, **L** is the **Gaussian-smoothed** image at its selected scale,

$$m(x,y) = \sqrt{\underbrace{(L(x+1,y) - L(x-1,y))^2}_{\text{x-derivative}} + \underbrace{(L(x,y+1) - L(x,y-1))^2}_{\text{y-derivative}}}$$

$$\theta(x,y) = \tan^{-1}((L(x,y+1) - L(x,y-1))/(L(x+1,y) - L(x-1,y)))$$

Detection process returns $\{\underbrace{x, y}_{\text{location}}, \underbrace{\sigma}_{\text{scale}}, \underbrace{\theta}_{\text{orientation}}\}$

4. Keypoint descriptor

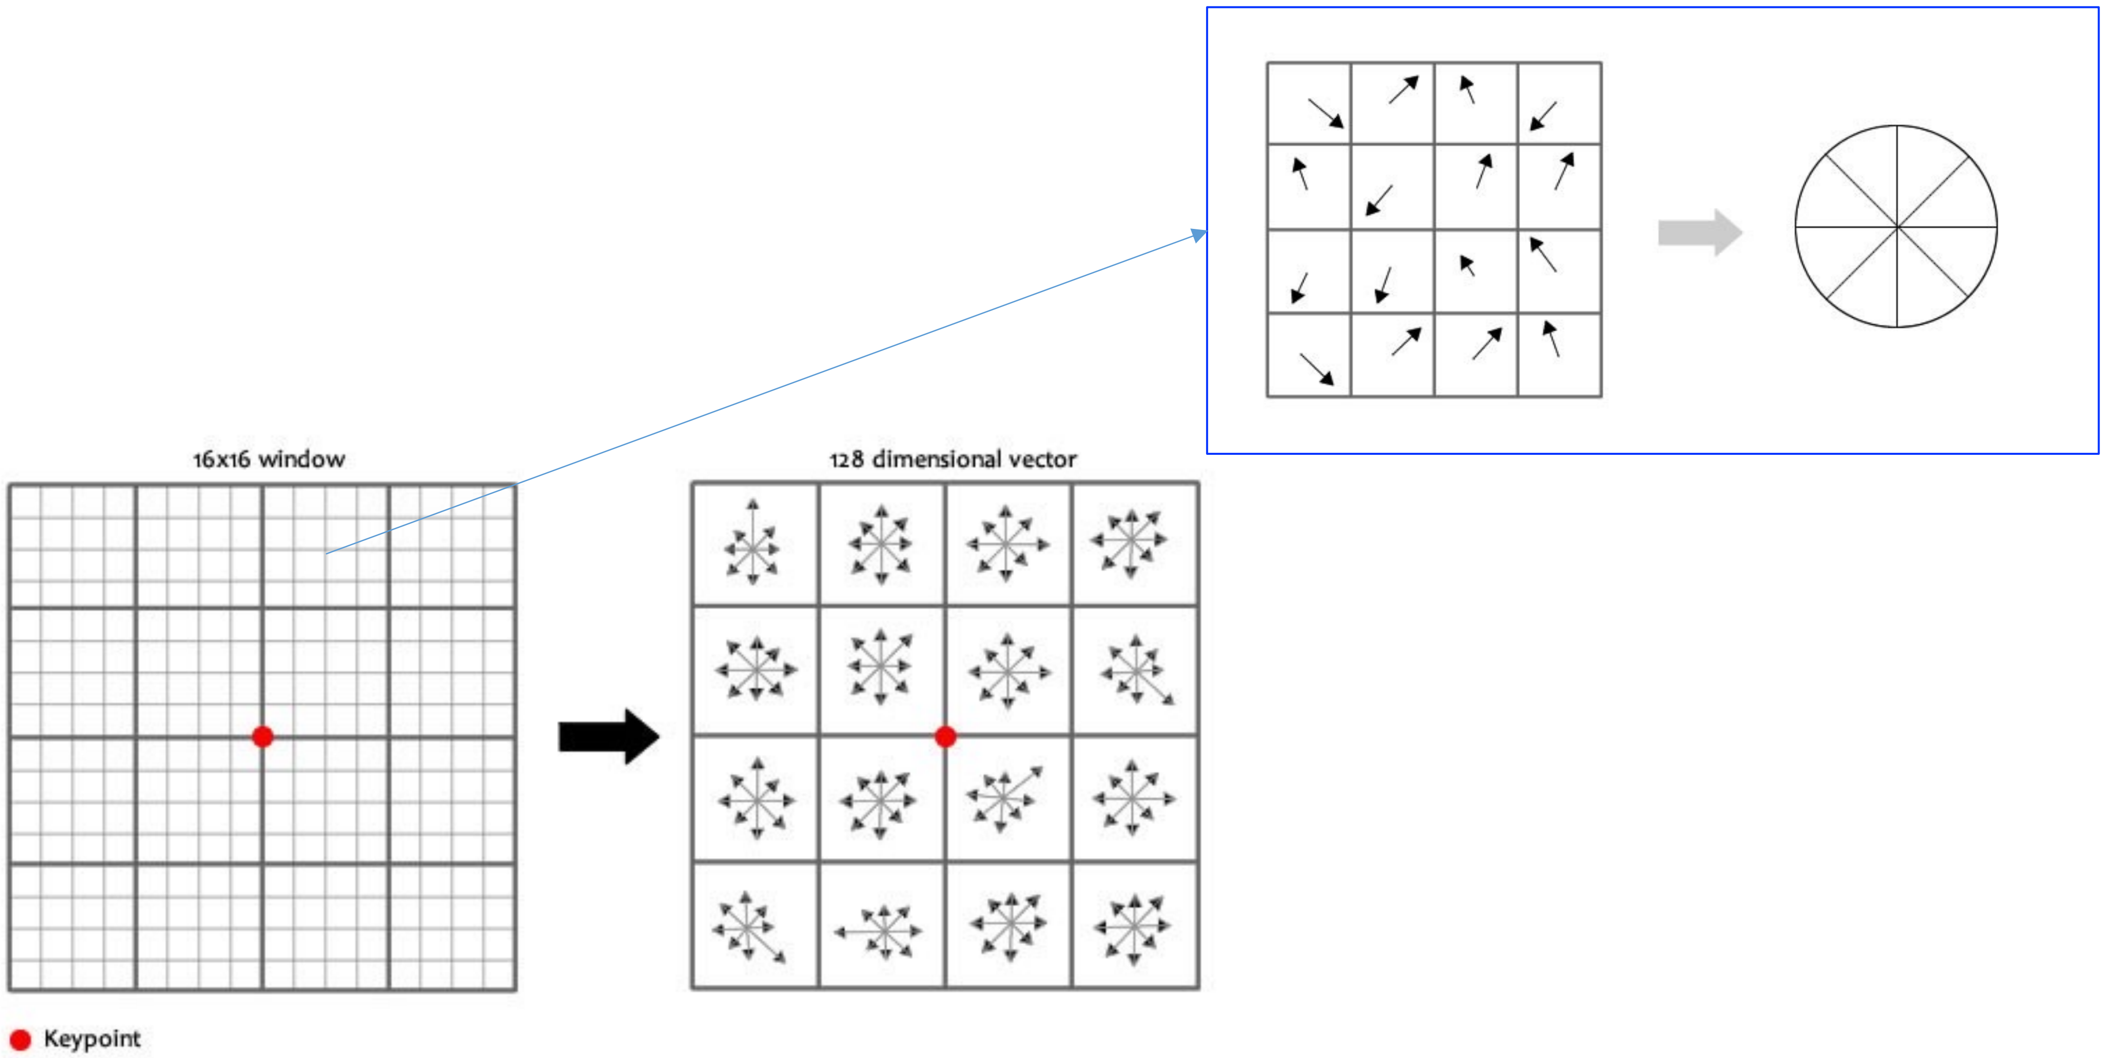

Adding more invariances to SIFT

Rotation Invariance:

- The feature vector uses gradient orientations. So if you rotate the image, everything changes!
- SIFT adopts “relative rotation”: the keypoint’s rotation is subtracted from each orientation. Thus each gradient orientation is relative to the keypoint’s orientation.

Illumination Invariance:

- All keypoints’ 128-dim vector are normalized to 1
- Sometimes we have ”outlier illumination” …
    - Practically, after normalization, we clamp all gradients > 0.2, then renormalize to [0,1]

*SIFT achieves an extremely elegant and robust balance between **global layout** (histogram) versus **local feature** (full gradient), discriminativeness versus resilience*


inliers: 687/734
TEXTS IN COMPUTER SCIENCE
Computer Vision
Algorithms and Applications
Richard Szeliski
Springer

Review: Local Descriptors

- Most features can be thought of as templates, histograms (counts), or combinations
- The ideal descriptor should be
  - Robust and Distinctive
  - Compact and Efficient
- Most available descriptors focus on edge/gradient information
  - Capture texture information
  - Color rarely used

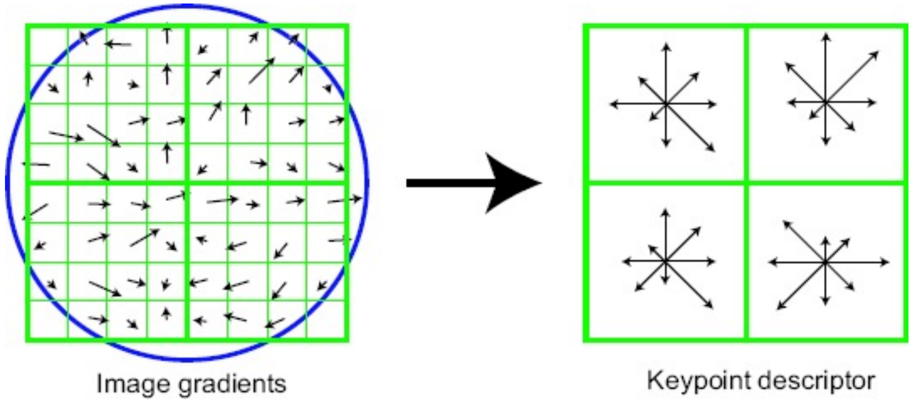

K. Grauman, B. Leibe

Local features: main components

1) **Detection:**
Find a set of distinctive key points.

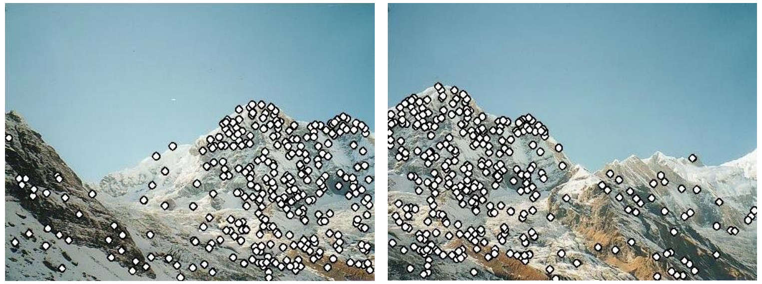

2) **Description:**
Extract feature descriptor around each interest point as vector.

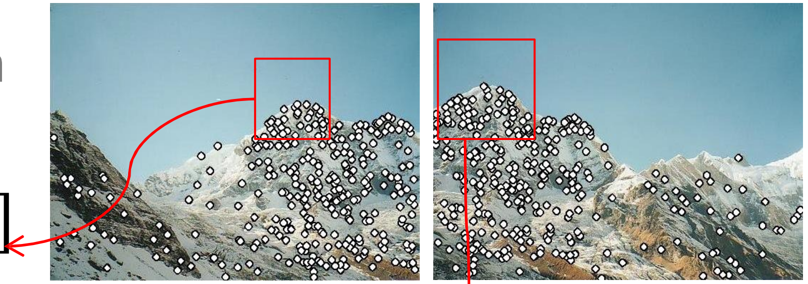

x_1

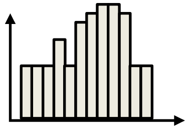

$$\mathbf{x}_1 = [x_1^{(1)}, \dots, x_d^{(1)}]$$

x_2

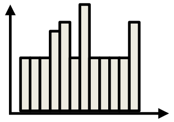

$$\mathbf{x}_2 = [x_1^{(2)}, \dots, x_d^{(2)}]$$

3) **Matching:**
Compute distance between feature vectors to find correspondence.

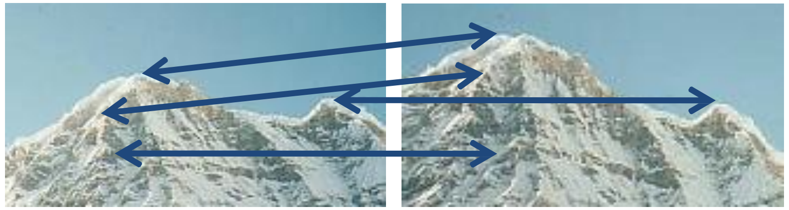

K. Grauman, B. Leibe

Ok, now we have local features…

But how similar can the two features be called “match”?

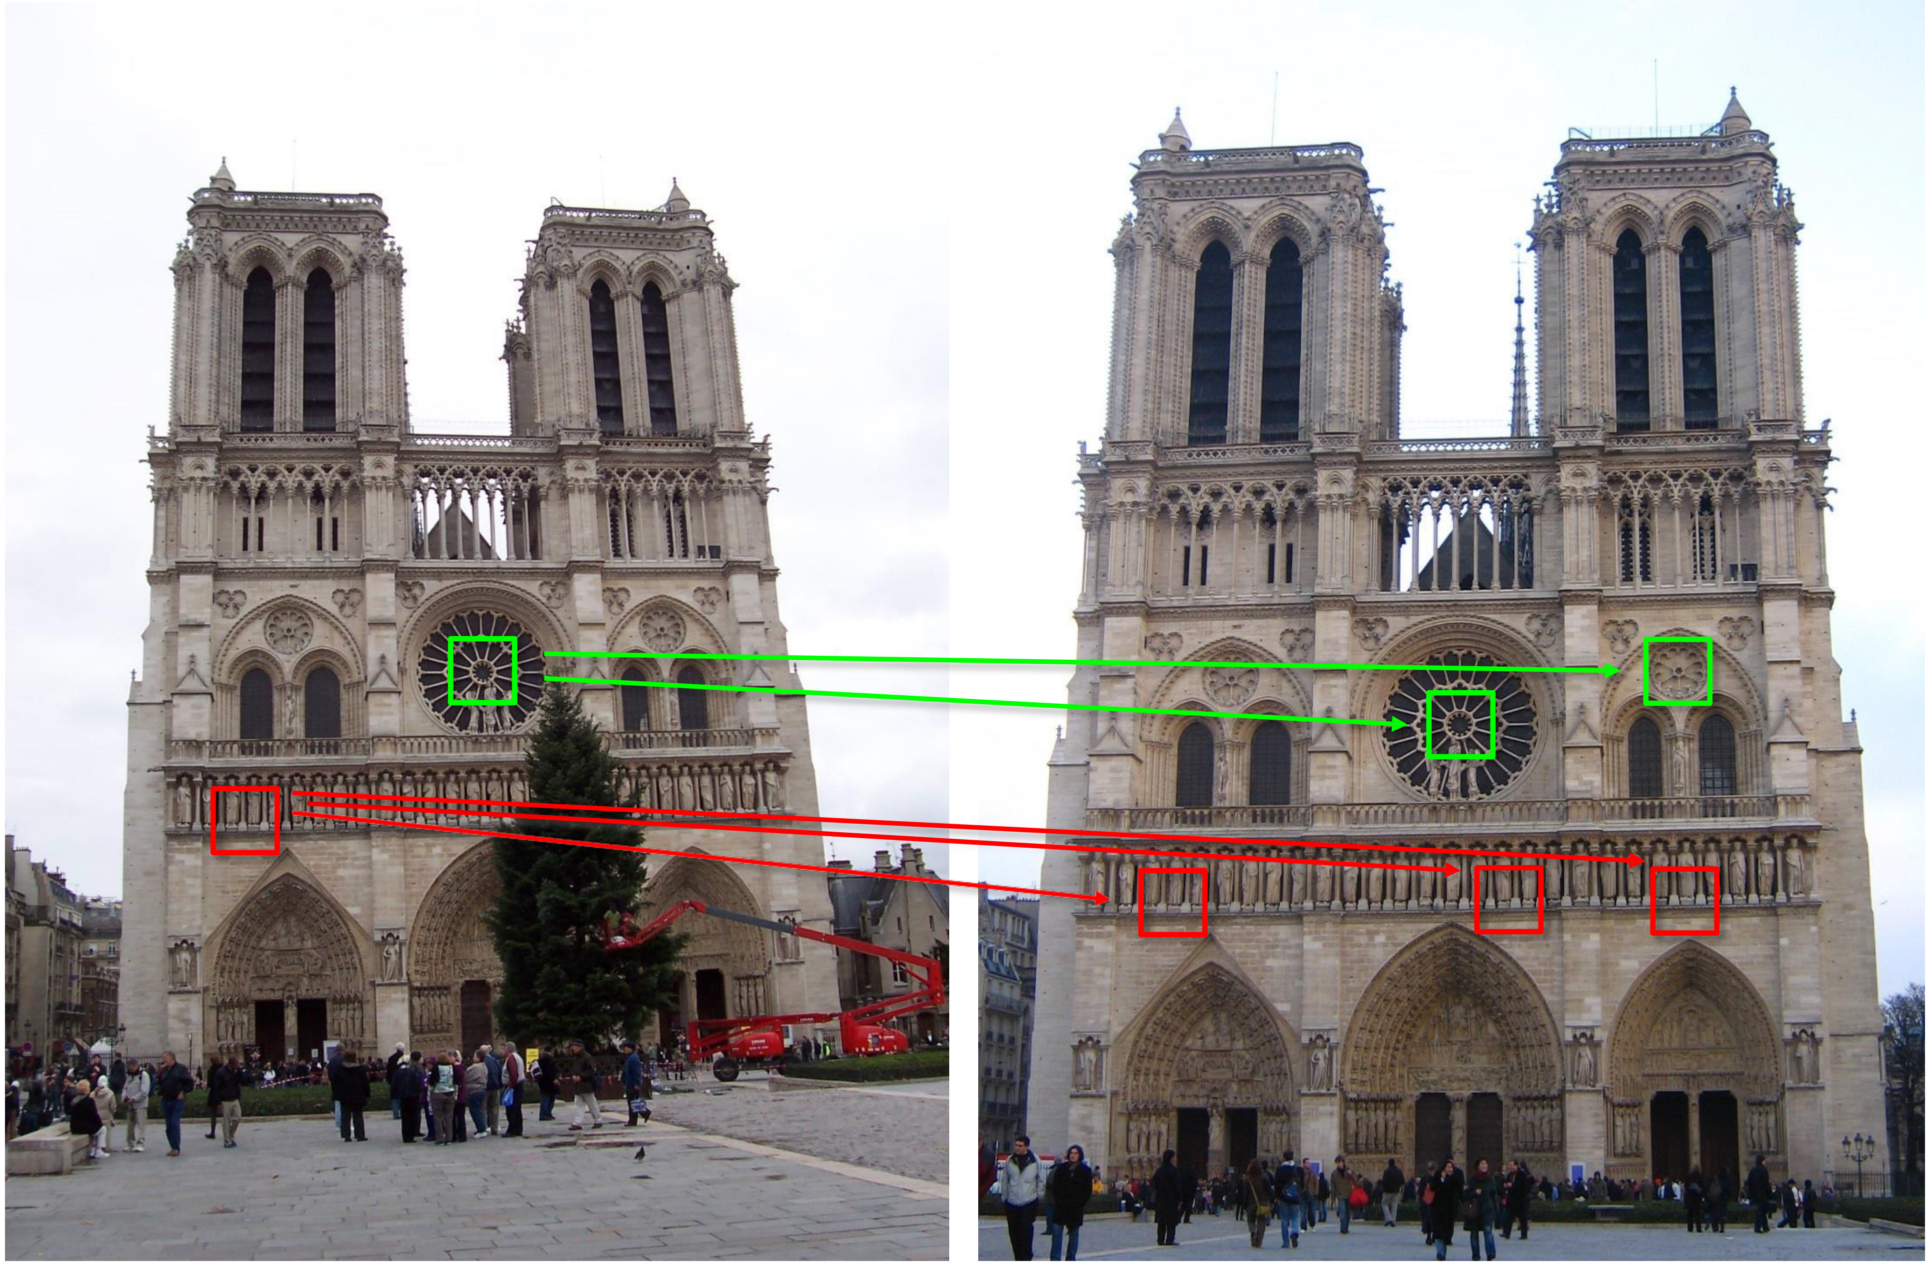

Distance: 0.34, 0.30, 0.40

Distance: 0.61, 1.22

What to Consider in the Design of Feature Matching

- Two images, I_1 and I_2

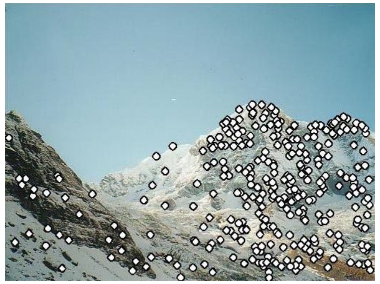

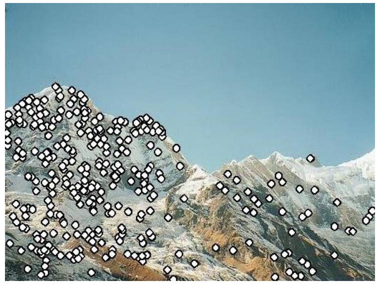

- Two sets X_1 and X_2 of feature points
  - Each feature point **x₁** has a descriptor $\mathbf{x}_1 = [x_1^{(1)}, \ldots, x_d^{(1)}]$

- Distance, bijective/injective/surjective, noise, confidence, computational complexity, generality…

Euclidean distance vs. Cosine Similarity

- Euclidean distance:

$$\mathrm{d}(\mathbf{p},\mathbf{q}) = \mathrm{d}(\mathbf{q},\mathbf{p}) = \sqrt{(q_1-p_1)^2 + (q_2-p_2)^2 + \cdots + (q_n-p_n)^2}$$

$$= \sqrt{\sum_{i=1}^{n}(q_i-p_i)^2}.$$

$$\|\mathbf{q}-\mathbf{p}\| = \sqrt{(\mathbf{q}-\mathbf{p})\cdot(\mathbf{q}-\mathbf{p})}.$$

- Cosine similarity:

$$\mathbf{a}\cdot\mathbf{b} = \|\mathbf{a}\|_2 \|\mathbf{b}\|_2 \cos\theta$$

$$\text{similarity} = \cos(\theta) = \frac{\mathbf{A}\cdot\mathbf{B}}{\|\mathbf{A}\|_2\|\mathbf{B}\|_2}$$

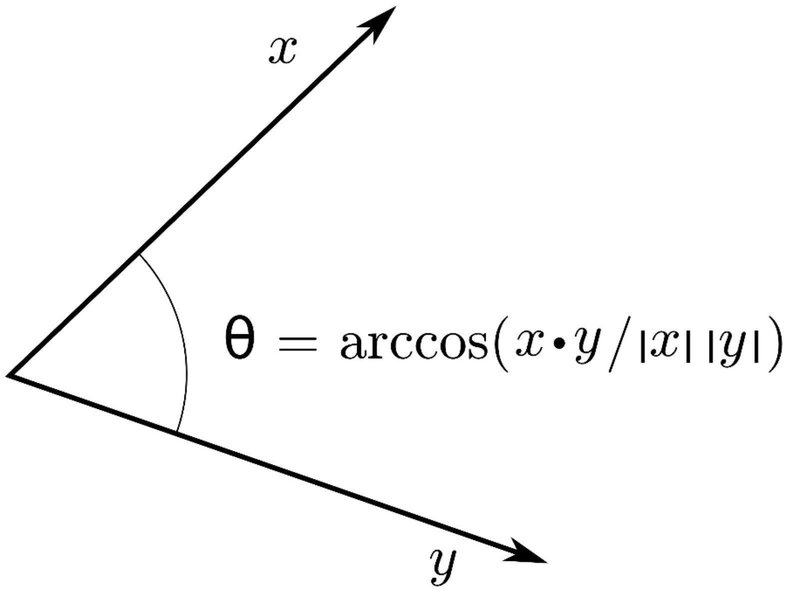

Wikipedia

Feature Matching

- Criteria 1:
  - Compute distance in feature space, e.g., Euclidean distance between 128-dim SIFT descriptors
  - Match point to lowest distance (nearest neighbor)

- Problems:
  - Does everything have a match?

Feature Matching

- Criteria 2:
  - Compute distance in feature space, e.g., Euclidean distance between 128-dim SIFT descriptors
  - Match point to lowest distance (nearest neighbor)
  - Ignore anything higher than threshold (no match!)

- Problems:
  - Threshold is hard to pick
  - Non-distinctive features could have lots of close matches, only one of which is correct

Nearest Neighbor Distance Ratio

Compare distance of closest (NN1) and second-closest (NN2) feature vector neighbor.

- If NN1 ≈ NN2, ratio $\frac{NN1}{NN2}$ will be ≈ 1 -> matches too close.
- As NN1 << NN2, ratio $\frac{NN1}{NN2}$ tends to 0.

Sorting by this ratio puts matches in order of confidence.

Threshold ratio – but how to choose?

Nearest Neighbor Distance Ratio

- Lowe computed a probability distribution functions of ratios
- 40,000 keypoints with hand-labeled ground truth

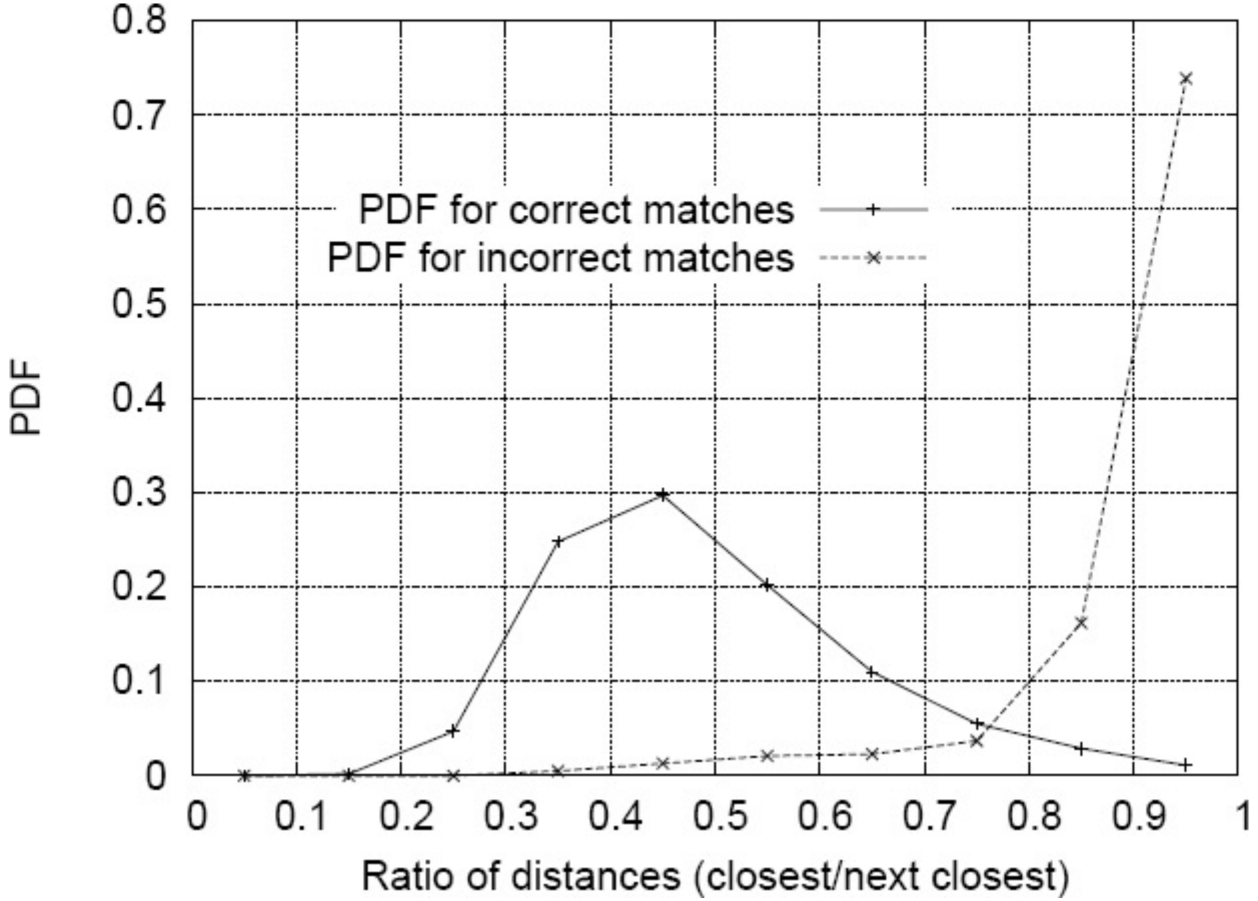

Ratio threshold depends on your application's view on the trade-off between the number of false positives and true positives!

Lowe IJCV 2004

Visual Similarity is still/forever an OPEN problem

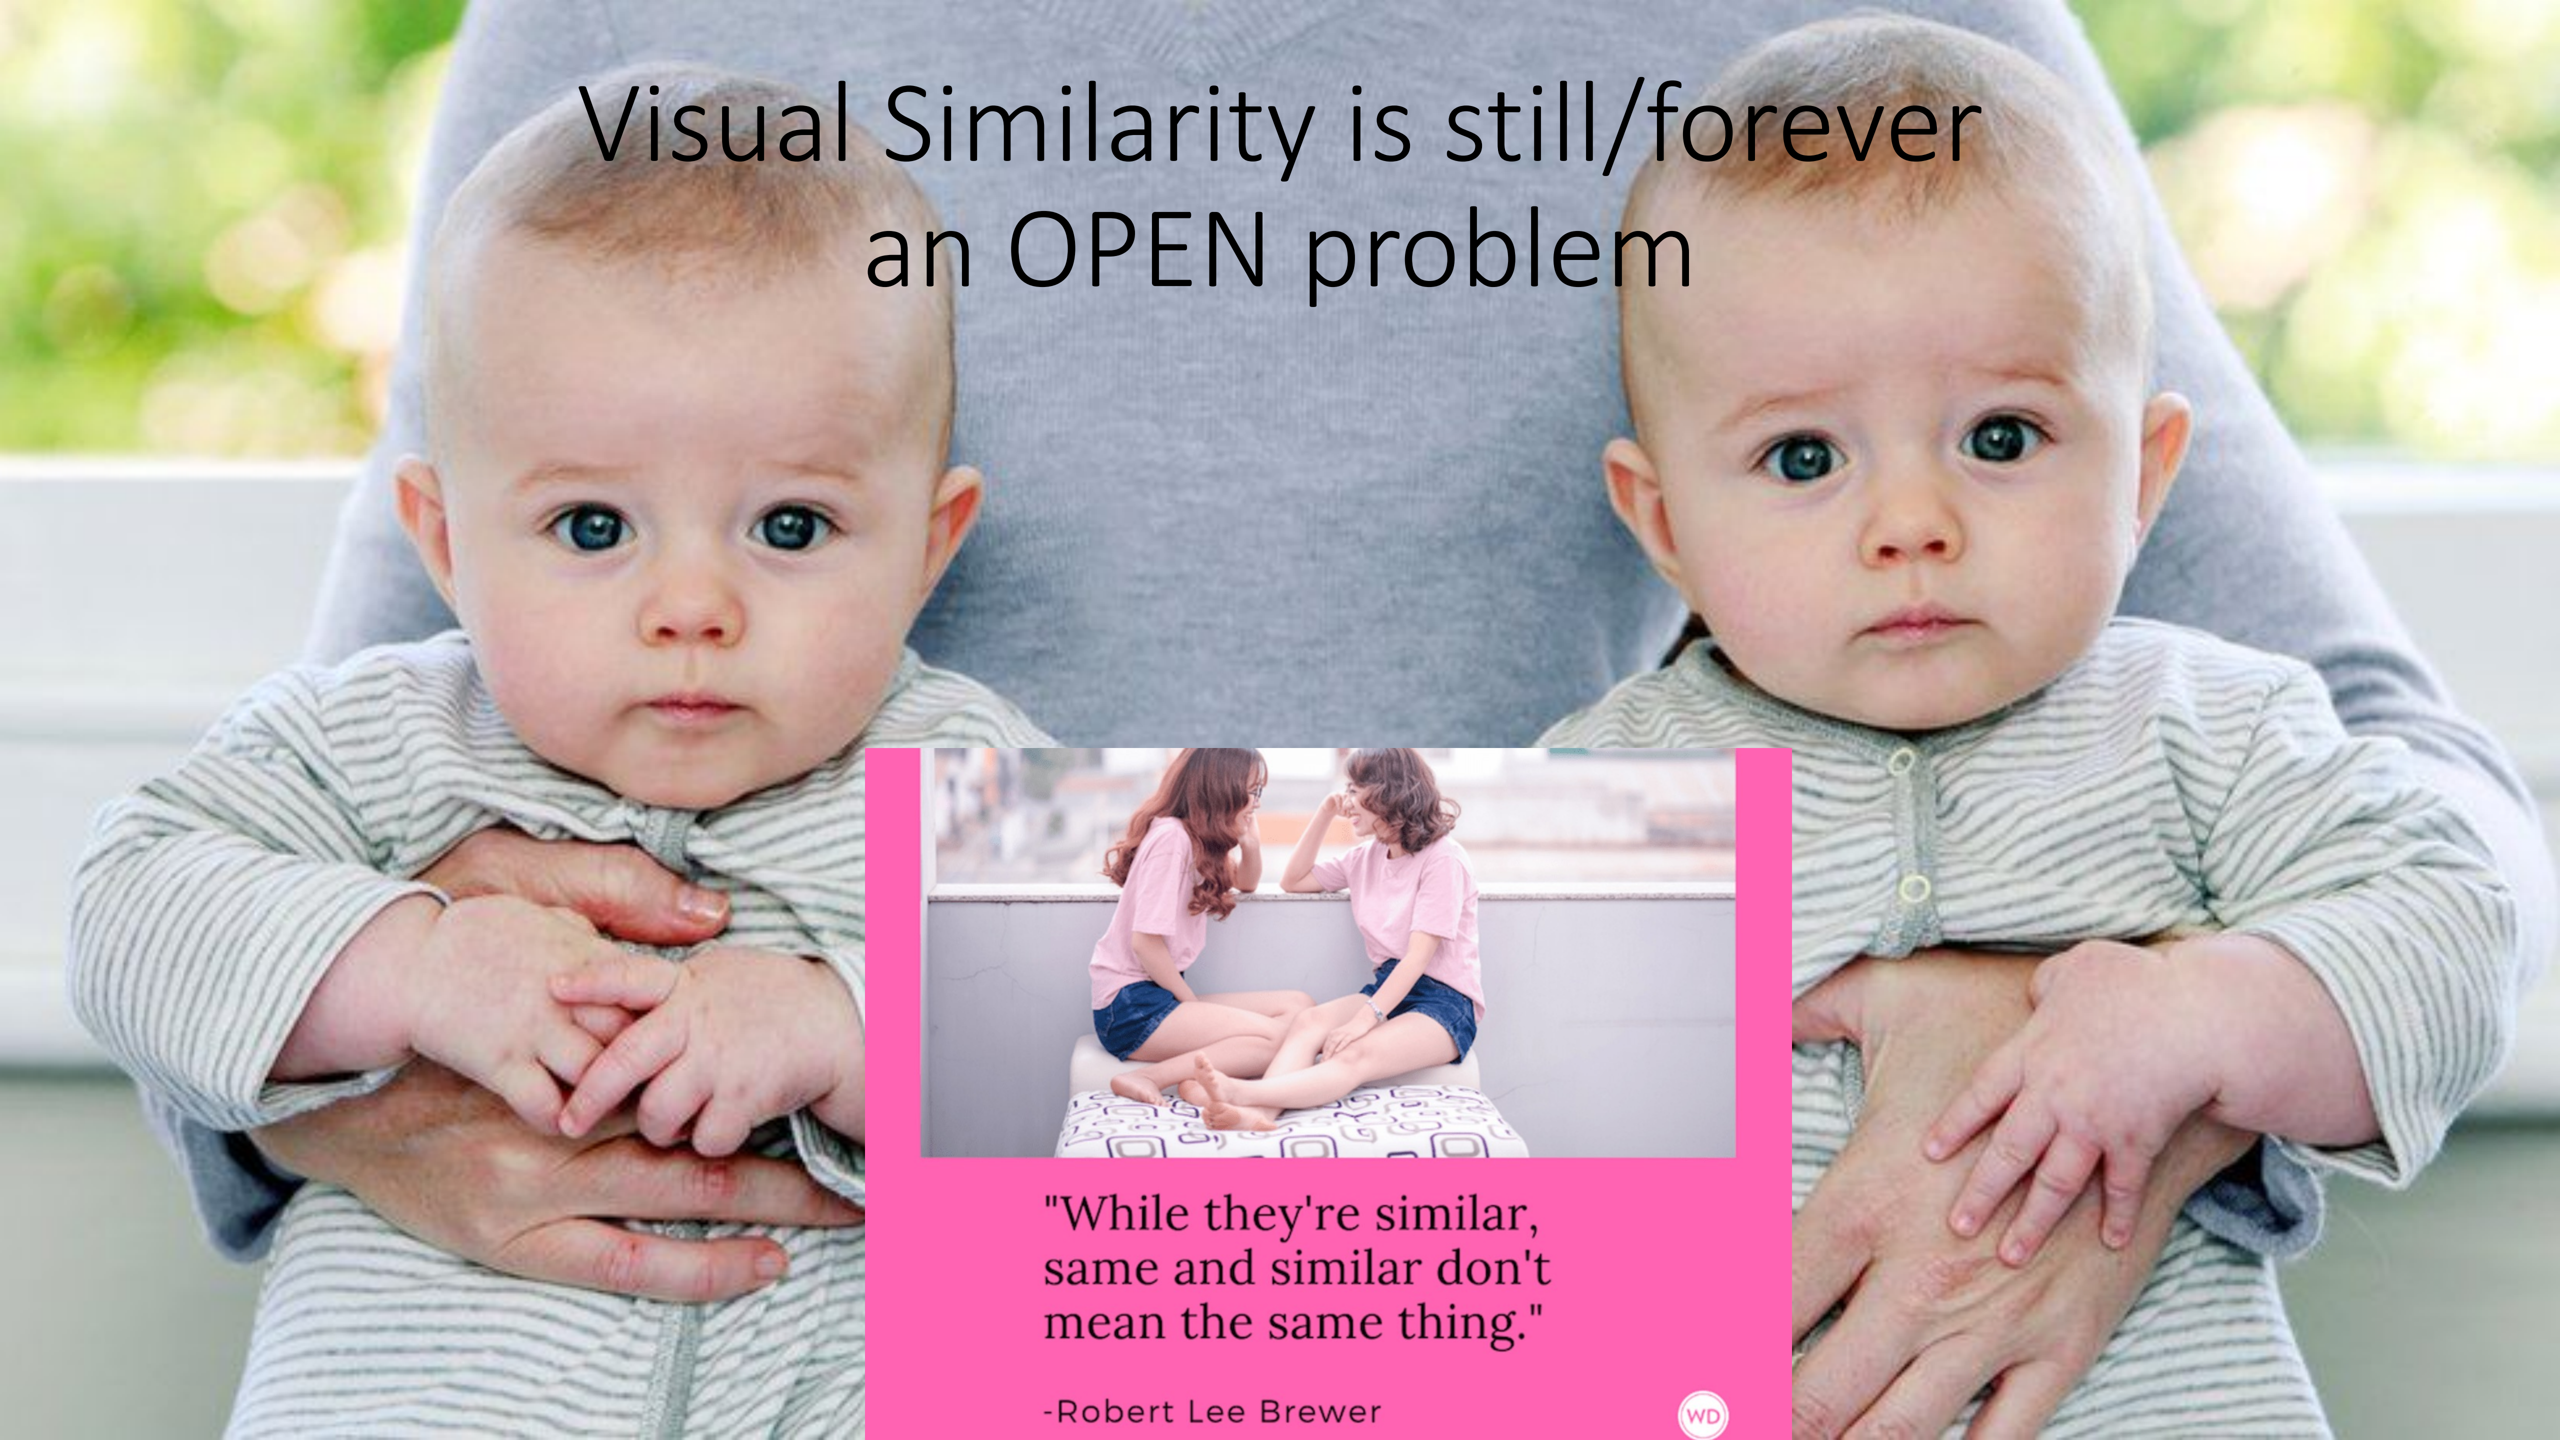

The University of Texas at Austin
Electrical and Computer Engineering
Cockrell School of Engineering